

Data Loading

Importing of libraries

```
In [2]: # Import necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Load datasets using the absolute paths

```
In [3]: import pandas as pd

# Load datasets
crashes = pd.read_csv('C:/Users/MNJORGE16/Desktop/Moringa/phase_3/project_
people = pd.read_csv('C:/Users/MNJORGE16/Desktop/Moringa/phase_3/project_
vehicles = pd.read_csv('C:/Users/MNJORGE16/Desktop/Moringa/phase_3/project_
```

c:\Users\MNJORGE16\AppData\Local\anaconda3\envs\learn-env\lib\site-packages\IPython\core\interactiveshell.py:3145: DtypeWarning: Columns (19,23,24,25,28) have mixed types.Specify dtype option on import or set low_memory=False.

has_raised = await self.run_ast_nodes(code_ast.body, cell_name, c:\Users\MNJORGE16\AppData\Local\anaconda3\envs\learn-env\lib\site-packages\IPython\core\interactiveshell.py:3145: DtypeWarning: Columns (20,39,40,41,43,47,48,49,52,54,57,58,60,70) have mixed types.Specify dtype option on import or set low_memory=False.

has_raised = await self.run_ast_nodes(code_ast.body, cell_name,

Data Inspection

Reviewing each dataset structure (rows, columns)

```
In [4]: print(crashes.shape)
```

(866411, 48)

```
In [5]: print(people.shape)
```

(1467049, 29)

```
In [6]: print(vehicles.shape)
```

```
(1764900, 71)
```

Inspecting the first few rows of each data set

This is to confirm data is loaded correctly

```
In [7]: print(crashes.head())
```

```
CRASH_RECORD_ID CRASH_DATE_EST_I \
0 23a79931ef555d54118f64dc9be2cf2dbf59636ce253f7... NaN
1 2675c13fd0f474d730a5b780968b3cafc7c12d7adb661f... NaN
2 5f54a59fcb087b12ae5b1acff96a3caf4f2d37e79f8db4... NaN
3 7ebf015016f83d09b321afd671a836d6b148330535d5df... NaN
4 6c1659069e9c6285a650e70d6f9b574ed5f64c12888479... NaN

CRASH_DATE POSTED_SPEED_LIMIT TRAFFIC_CONTROL_DEVICE \
0 09/05/2023 07:05:00 PM 30 TRAFFIC SIGNAL
1 09/22/2023 06:45:00 PM 50 NO CONTROLS
2 07/29/2023 02:45:00 PM 30 TRAFFIC SIGNAL
3 08/09/2023 11:00:00 PM 30 NO CONTROLS
4 08/18/2023 12:50:00 PM 15 OTHER

DEVICE_CONDITION WEATHER_CONDITION LIGHTING_CONDITION \
0 FUNCTIONING PROPERLY CLEAR DUSK
1 NO CONTROLS CLEAR DARKNESS, LIGHTED ROAD
2 FUNCTIONING PROPERLY CLEAR DAYLIGHT
3 NO CONTROLS CLEAR DARKNESS, LIGHTED ROAD
4 FUNCTIONING PROPERLY CLEAR DAYLIGHT

FIRST_CRASH_TYPE TRAFFICWAY_TYPE ... \
0 ANGLE FIVE POINT, OR MORE ...
1 REAR END DIVIDED - W/MEDIAN BARRIER ...
2 PARKED MOTOR VEHICLE DIVIDED - W/MEDIAN (NOT RAISED) ...
3 SIDESWIPE SAME DIRECTION NOT DIVIDED ...
4 REAR END OTHER ...

INJURIES_NON_INCAPACITATING INJURIES_REPORTED_NOT_EVIDENT \
0 2.0 0.0
1 0.0 0.0
2 0.0 0.0
3 0.0 0.0
4 1.0 0.0

INJURIES_NO_INDICATION INJURIES_UNKNOWN CRASH_HOUR CRASH_DAY_OF_WEEK \
0 2.0 0.0 19 3
1 2.0 0.0 18 6
2 1.0 0.0 14 7
3 2.0 0.0 23 4
4 1.0 0.0 12 6

CRASH_MONTH LATITUDE LONGITUDE LOCAT
ION
0 9 NaN NaN
NaN
1 9 NaN NaN
NaN
2 7 41.85412 -87.665902 POINT (-87.665902342962 41.8541202629
52)
3 8 NaN NaN
NaN
4 8 NaN NaN
NaN
```

```
[5 rows x 48 columns]
```

```
In [8]: print(people.head())
```

```
PERSON_ID PERSON_TYPE CRASH_RECORD_I
D \
0 0749947 DRIVER 81dc0de2ed92aa62baccab641fa377be7feb1cc47e655
4...
1 0871921 DRIVER af84fb5c8d996fcd3aefd36593c3a02e6e7509eeb2756
8...
2 010018 DRIVER 71162af7bf22799b776547132ebf134b5b438dcf3dac6
b...
3 010038 DRIVER c21c476e2ccc41af550b5d858d22aaac4ffc88745a170
0...
4 010039 DRIVER eb390a4c8e114c69488f5fb8a097fe629f5a92fd528cf
4...

VEHICLE_ID CRASH_DATE SEAT_NO CITY STATE ZIPCODE SEX
\
0 834816.0 09/28/2019 03:30:00 AM NaN CHICAGO IL 60651 M
1 827212.0 04/13/2020 10:50:00 PM NaN CHICAGO IL 60620 M
2 9579.0 11/01/2015 05:00:00 AM NaN NaN NaN NaN X
3 9598.0 11/01/2015 08:00:00 AM NaN NaN NaN NaN X
4 9600.0 11/01/2015 10:15:00 AM NaN NaN NaN NaN X

... EMS_RUN_NO DRIVER_ACTION DRIVER_VISION PHYSICAL_CONDITION \
0 ... NaN UNKNOWN UNKNOWN UNKNOWN
1 ... NaN NONE NOT OBSCURED NORMAL
2 ... NaN IMPROPER BACKING UNKNOWN UNKNOWN
3 ... NaN UNKNOWN UNKNOWN UNKNOWN
4 ... NaN UNKNOWN UNKNOWN UNKNOWN

PEDPEDAL_ACTION PEDPEDAL_VISIBILITY PEDPEDAL_LOCATION BAC_RESULT
\
0 NaN NaN NaN TEST NOT OFFERED
1 NaN NaN NaN TEST NOT OFFERED
2 NaN NaN NaN TEST NOT OFFERED
3 NaN NaN NaN TEST NOT OFFERED
4 NaN NaN NaN TEST NOT OFFERED

BAC_RESULT VALUE CELL_PHONE_USE
0 NaN NaN
1 NaN NaN
2 NaN NaN
3 NaN NaN
4 NaN NaN
```

```
[5 rows x 29 columns]
```

```
In [9]: print(vehicles.head())
```

```
CRASH_UNIT_ID          CRASH_RECORD_ID \
0      1727162  f5943b05f46b8d4148a63b7506a59113eae0cf1075aabc...
1      1717556  7b1763088507f77e0e552c009a6bf89a4d6330c7527706...
2      1717574  2603ff5a88f0b9b54576934c5ed4e4a64e8278e005687b...
3      1717579  a52ef70e33d468b855b5be44e8638a564434dcf99c0edf...
4      1720118  609055f4b1a72a44d6ec40ba9036cefd7c1287a755eb6c...

      CRASH_DATE  UNIT_NO  UNIT_TYPE  NUM_PASSENGERS  VEHICLE_I
D \
0  12/21/2023 08:57:00 AM          2  PEDESTRIAN          NaN          Na
N
1  12/06/2023 03:24:00 PM          1    DRIVER          NaN  1634931.
0
2  12/06/2023 04:00:00 PM          2    DRIVER          NaN  1634978.
0
3  12/06/2023 04:30:00 PM          1    DRIVER          NaN  1634948.
0
4  12/10/2023 12:12:00 PM          1    DRIVER          NaN  1637401.
0

CMRC_VEH_I    MAKE    MODEL  ... TRAILER1_LENGTH  TRAILER2_LENGTH \
0          NaN      NaN      ...          NaN          NaN
1          NaN    NISSAN    SENTRA  ...          NaN          NaN
2          NaN  CHRYSLER    SEBRING  ...          NaN          NaN
3          NaN    SUBARU    OUTBACK  ...          NaN          NaN
4          NaN    TOYOTA    RAV4    ...          NaN          NaN

TOTAL_VEHICLE_LENGTH  AXLE_CNT  VEHICLE_CONFIG  CARGO_BODY_TYPE  LOAD_TYPE
\
0          NaN          NaN          NaN          NaN          NaN
1          NaN          NaN          NaN          NaN          NaN
2          NaN          NaN          NaN          NaN          NaN
3          NaN          NaN          NaN          NaN          NaN
4          NaN          NaN          NaN          NaN          NaN

HAZMAT_OUT_OF_SERVICE_I  MCS_OUT_OF_SERVICE_I  HAZMAT_CLASS
0          NaN          NaN          NaN
1          NaN          NaN          NaN
2          NaN          NaN          NaN
3          NaN          NaN          NaN
4          NaN          NaN          NaN
```

```
[5 rows x 71 columns]
```

Inspecting Column Names

```
In [10]: print(crashes.columns)
```

```
Index(['CRASH_RECORD_ID', 'CRASH_DATE_EST_I', 'CRASH_DATE',  
      'POSTED_SPEED_LIMIT', 'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDI  
TION',  
      'WEATHER_CONDITION', 'LIGHTING_CONDITION', 'FIRST_CRASH_TYPE',  
      'TRAFFICWAY_TYPE', 'LANE_CNT', 'ALIGNMENT', 'ROADWAY_SURFACE_CON  
D',  
      'ROAD_DEFECT', 'REPORT_TYPE', 'CRASH_TYPE', 'INTERSECTION_RELATED_  
I',  
      'NOT_RIGHT_OF_WAY_I', 'HIT_AND_RUN_I', 'DAMAGE', 'DATE_POLICE_NOTI  
FIED',  
      'PRIM_CONTRIBUTORY_CAUSE', 'SEC_CONTRIBUTORY_CAUSE', 'STREET_NO',  
      'STREET_DIRECTION', 'STREET_NAME', 'BEAT_OF_OCCURRENCE',  
      'PHOTOS_TAKEN_I', 'STATEMENTS_TAKEN_I', 'DOORING_I', 'WORK_ZONE_  
I',  
      'WORK_ZONE_TYPE', 'WORKERS_PRESENT_I', 'NUM_UNITS',  
      'MOST_SEVERE_INJURY', 'INJURIES_TOTAL', 'INJURIES_FATAL',  
      'INJURIES_INCAPACITATING', 'INJURIES_NON_INCAPACITATING',  
      'INJURIES_REPORTED_NOT_EVIDENT', 'INJURIES_NO_INDICATION',  
      'INJURIES_UNKNOWN', 'CRASH_HOUR', 'CRASH_DAY_OF_WEEK', 'CRASH_MONT  
H',  
      'LATITUDE', 'LONGITUDE', 'LOCATION'],  
      dtype='object')
```

```
In [11]: print(people.columns)
```

```
Index(['PERSON_ID', 'PERSON_TYPE', 'CRASH_RECORD_ID', 'VEHICLE_ID',  
      'CRASH_DATE', 'SEAT_NO', 'CITY', 'STATE', 'ZIPCODE', 'SEX', 'AGE',  
      'DRIVERS_LICENSE_STATE', 'DRIVERS_LICENSE_CLASS', 'SAFETY_EQUIPMEN  
T',  
      'AIRBAG_DEPLOYED', 'EJECTION', 'INJURY_CLASSIFICATION', 'HOSPITA  
L',  
      'EMS_AGENCY', 'EMS_RUN_NO', 'DRIVER_ACTION', 'DRIVER_VISION',  
      'PHYSICAL_CONDITION', 'PEDPEDAL_ACTION', 'PEDPEDAL_VISIBILITY',  
      'PEDPEDAL_LOCATION', 'BAC_RESULT', 'BAC_RESULT VALUE',  
      'CELL_PHONE_USE'],  
      dtype='object')
```

```
In [12]: print(vehicles.columns)
```

```
Index(['CRASH_UNIT_ID', 'CRASH_RECORD_ID', 'CRASH_DATE', 'UNIT_NO',  
      'UNIT_TYPE', 'NUM_PASSENGERS', 'VEHICLE_ID', 'CMRC_VEH_I', 'MAKE',  
      'MODEL', 'LIC_PLATE_STATE', 'VEHICLE_YEAR', 'VEHICLE_DEFECT',  
      'VEHICLE_TYPE', 'VEHICLE_USE', 'TRAVEL_DIRECTION', 'MANEUVER',  
      'TOWED_I', 'FIRE_I', 'OCCUPANT_CNT', 'EXCEED_SPEED_LIMIT_I', 'TOWE  
D_BY',  
      'TOWED_TO', 'AREA_00_I', 'AREA_01_I', 'AREA_02_I', 'AREA_03_I',  
      'AREA_04_I', 'AREA_05_I', 'AREA_06_I', 'AREA_07_I', 'AREA_08_I',  
      'AREA_09_I', 'AREA_10_I', 'AREA_11_I', 'AREA_12_I', 'AREA_99_I',  
      'FIRST_CONTACT_POINT', 'CMV_ID', 'USDOT_NO', 'CCMC_NO', 'ILCC_NO',  
      'COMMERCIAL_SRC', 'GVWR', 'CARRIER_NAME', 'CARRIER_STATE',  
      'CARRIER_CITY', 'HAZMAT_PLACARDS_I', 'HAZMAT_NAME', 'UN_NO',  
      'HAZMAT_PRESENT_I', 'HAZMAT_REPORT_I', 'HAZMAT_REPORT_NO',  
      'MCS_REPORT_I', 'MCS_REPORT_NO', 'HAZMAT_VIO_CAUSE_CRASH_I',  
      'MCS_VIO_CAUSE_CRASH_I', 'IDOT_PERMIT_NO', 'WIDE_LOAD_I',  
      'TRAILER1_WIDTH', 'TRAILER2_WIDTH', 'TRAILER1_LENGTH',  
      'TRAILER2_LENGTH', 'TOTAL_VEHICLE_LENGTH', 'AXLE_CNT', 'VEHICLE_CO  
NFIG',  
      'CARGO_BODY_TYPE', 'LOAD_TYPE', 'HAZMAT_OUT_OF_SERVICE_I',  
      'MCS_OUT_OF_SERVICE_I', 'HAZMAT_CLASS'],  
      dtype='object')
```

Data Types and Schema

```
In [13]: # Check data types of each dataset
print(crashes.dtypes)

# Identify any incorrect data types (e.g., numeric columns read as objects)
```

CRASH_RECORD_ID	object
CRASH_DATE_EST_I	object
CRASH_DATE	object
POSTED_SPEED_LIMIT	int64
TRAFFIC_CONTROL_DEVICE	object
DEVICE_CONDITION	object
WEATHER_CONDITION	object
LIGHTING_CONDITION	object
FIRST_CRASH_TYPE	object
TRAFFICWAY_TYPE	object
LANE_CNT	float64
ALIGNMENT	object
ROADWAY_SURFACE_COND	object
ROAD_DEFECT	object
REPORT_TYPE	object
CRASH_TYPE	object
INTERSECTION_RELATED_I	object
NOT_RIGHT_OF_WAY_I	object
HIT_AND_RUN_I	object
DAMAGE	object
DATE_POLICE_NOTIFIED	object
PRIM_CONTRIBUTORY_CAUSE	object
SEC_CONTRIBUTORY_CAUSE	object
STREET_NO	int64
STREET_DIRECTION	object
STREET_NAME	object
BEAT_OF_OCCURRENCE	float64
PHOTOS_TAKEN_I	object
STATEMENTS_TAKEN_I	object
DOORING_I	object
WORK_ZONE_I	object
WORK_ZONE_TYPE	object
WORKERS_PRESENT_I	object
NUM_UNITS	int64
MOST_SEVERE_INJURY	object
INJURIES_TOTAL	float64
INJURIES_FATAL	float64
INJURIES_INCAPACITATING	float64
INJURIES_NON_INCAPACITATING	float64
INJURIES_REPORTED_NOT_EVIDENT	float64
INJURIES_NO_INDICATION	float64
INJURIES_UNKNOWN	float64
CRASH_HOUR	int64
CRASH_DAY_OF_WEEK	int64
CRASH_MONTH	int64
LATITUDE	float64
LONGITUDE	float64
LOCATION	object
dtype:	object


```
In [14]: print(people.dtypes)
```

```
PERSON_ID          object
PERSON_TYPE        object
CRASH_RECORD_ID    object
VEHICLE_ID         float64
CRASH_DATE         object
SEAT_NO           float64
CITY              object
STATE            object
ZIPCODE          object
SEX              object
AGE             float64
DRIVERS_LICENSE_STATE  object
DRIVERS_LICENSE_CLASS  object
SAFETY_EQUIPMENT    object
AIRBAG_DEPLOYED    object
EJECTION           object
INJURY_CLASSIFICATION  object
HOSPITAL           object
EMS_AGENCY         object
EMS_RUN_NO         object
DRIVER_ACTION      object
DRIVER_VISION      object
PHYSICAL_CONDITION object
PEDPEDAL_ACTION    object
PEDPEDAL_VISIBILITY object
PEDPEDAL_LOCATION  object
BAC_RESULT         object
BAC_RESULT VALUE   float64
CELL_PHONE_USE     object
dtype: object
```

```
In [15]: print(vehicles.dtypes)
```

```
CRASH_UNIT_ID      int64
CRASH_RECORD_ID    object
CRASH_DATE         object
UNIT_NO           int64
UNIT_TYPE          object
...
CARGO_BODY_TYPE    object
LOAD_TYPE          object
HAZMAT_OUT_OF_SERVICE_I  object
MCS_OUT_OF_SERVICE_I  object
HAZMAT_CLASS       object
Length: 71, dtype: object
```

Initial Summary Statistics

```
In [16]: # Summary statistics for numerical features - Crashes  
print(crashes.describe())
```

	POSTED_SPEED_LIMIT	LANE_CNT	STREET_NO	BEAT_OF_OCCURENC
E \				
count	866411.000000	1.990150e+05	866411.000000	866406.00000
0				
mean	28.415733	1.332981e+01	3687.152034	1244.46922
7				
std	6.131785	2.961557e+03	2882.599171	705.12605
9				
min	0.000000	0.000000e+00	0.000000	111.00000
0				
25%	30.000000	2.000000e+00	1250.000000	714.00000
0				
50%	30.000000	2.000000e+00	3201.000000	1212.00000
0				
75%	30.000000	4.000000e+00	5580.000000	1822.00000
0				
max	99.000000	1.191625e+06	451100.000000	6100.00000
0				

	NUM_UNITS	INJURIES_TOTAL	INJURIES_FATAL	INJURIES_INCAPACITA
TING \				
count	866411.000000	864508.000000	864508.000000	864508.00
0000				
mean	2.035117	0.192690	0.001194	0.01
9823				
std	0.452753	0.570222	0.037455	0.16
4843				
min	1.000000	0.000000	0.000000	0.00
0000				
25%	2.000000	0.000000	0.000000	0.00
0000				
50%	2.000000	0.000000	0.000000	0.00
0000				
75%	2.000000	0.000000	0.000000	0.00
0000				
max	18.000000	21.000000	4.000000	10.00
0000				

	INJURIES_NON_INCAPACITATING	INJURIES_REPORTED_NOT_EVIDENT \
count	864508.000000	864508.000000
mean	0.108248	0.063426
std	0.424294	0.323900
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	21.000000	15.000000

	INJURIES_NO_INDICATION	INJURIES_UNKNOWN	CRASH_HOUR \
count	864508.000000	864508.0	866411.000000
mean	2.001795	0.0	13.205135
std	1.157261	0.0	5.573549
min	0.000000	0.0	0.000000
25%	1.000000	0.0	9.000000
50%	2.000000	0.0	14.000000
75%	2.000000	0.0	17.000000
max	61.000000	0.0	23.000000

	CRASH_DAY_OF_WEEK	CRASH_MONTH	LATITUDE	LONGITUDE
count	866411.000000	866411.000000	860273.000000	860273.000000
mean	4.122962	6.606381	41.855078	-87.673657

std	1.981495	3.377482	0.333591	0.677645
min	1.000000	1.000000	0.000000	-87.936193
25%	2.000000	4.000000	41.782879	-87.721774
50%	4.000000	7.000000	41.874945	-87.674177
75%	6.000000	10.000000	41.924490	-87.633463
max	7.000000	12.000000	42.022780	0.000000

```
In [17]: # Summary statistics for numerical features - People
print(people.describe())
```

	VEHICLE_ID	SEAT_NO	AGE	BAC_RESULT	VALUE
count	1.438245e+06	299132.000000	1.040853e+06	1775.000000	
mean	6.905905e+05	4.160906	3.781694e+01	0.169448	
std	4.038528e+05	2.198771	1.710846e+01	0.102295	
min	2.000000e+00	1.000000	-1.770000e+02	0.000000	
25%	3.444260e+05	3.000000	2.500000e+01	0.120000	
50%	6.826620e+05	3.000000	3.500000e+01	0.170000	
75%	1.034161e+06	5.000000	5.000000e+01	0.220000	
max	1.801497e+06	12.000000	1.100000e+02	1.000000	

```
In [18]: # Summary statistics for numerical features - Vehicles
print(vehicles.describe())
```

	CRASH_UNIT_ID	UNIT_NO	NUM_PASSENGERS	VEHICLE_ID	\
count	1.764900e+06	1.764900e+06	261313.000000	1.724034e+06	
mean	9.438774e+05	3.705683e+00	1.470750	8.976615e+05	
std	5.463825e+05	2.843844e+03	1.055718	5.186095e+05	
min	2.000000e+00	0.000000e+00	1.000000	2.000000e+00	
25%	4.698658e+05	1.000000e+00	1.000000	4.489182e+05	
50%	9.450715e+05	2.000000e+00	1.000000	8.959025e+05	
75%	1.417380e+06	2.000000e+00	2.000000	1.346214e+06	
max	1.888827e+06	3.778035e+06	59.000000	1.799377e+06	

	VEHICLE_YEAR	OCCUPANT_CNT	CMV_ID	TRAILER1_LENGTH	\
count	1.448829e+06	1.724034e+06	17859.000000	2393.000000	
mean	2.014207e+03	1.079142e+00	9960.036564	48.511910	
std	1.385204e+02	7.815274e-01	5757.641443	20.695514	
min	1.900000e+03	0.000000e+00	1.000000	1.000000	
25%	2.007000e+03	1.000000e+00	4918.500000	45.000000	
50%	2.013000e+03	1.000000e+00	9988.000000	53.000000	
75%	2.017000e+03	1.000000e+00	14967.500000	53.000000	
max	9.999000e+03	9.900000e+01	19878.000000	740.000000	

	TRAILER2_LENGTH	TOTAL_VEHICLE_LENGTH	AXLE_CNT
count	70.000000	2918.000000	4396.000000
mean	44.271429	53.225497	9.619882
std	28.008240	31.291466	392.233256
min	1.000000	1.000000	1.000000
25%	24.250000	35.000000	2.000000
50%	50.000000	53.000000	3.000000
75%	53.000000	66.000000	5.000000
max	123.000000	999.000000	26009.000000

Data Relationships Before Dropping Columns with High Missing Values

Scatter Plots

To visualize relationships between pairs of numerical feature to spot trends, correlations, or anomalies.


Correlation Matrix

To understand how numerical features are related to one another, which is crucial for avoiding multicollinearity and for feature selection.


Cross-Tabulation

To explore relationships between categorical variables to show how the distribution of one categorical variable is related to another.

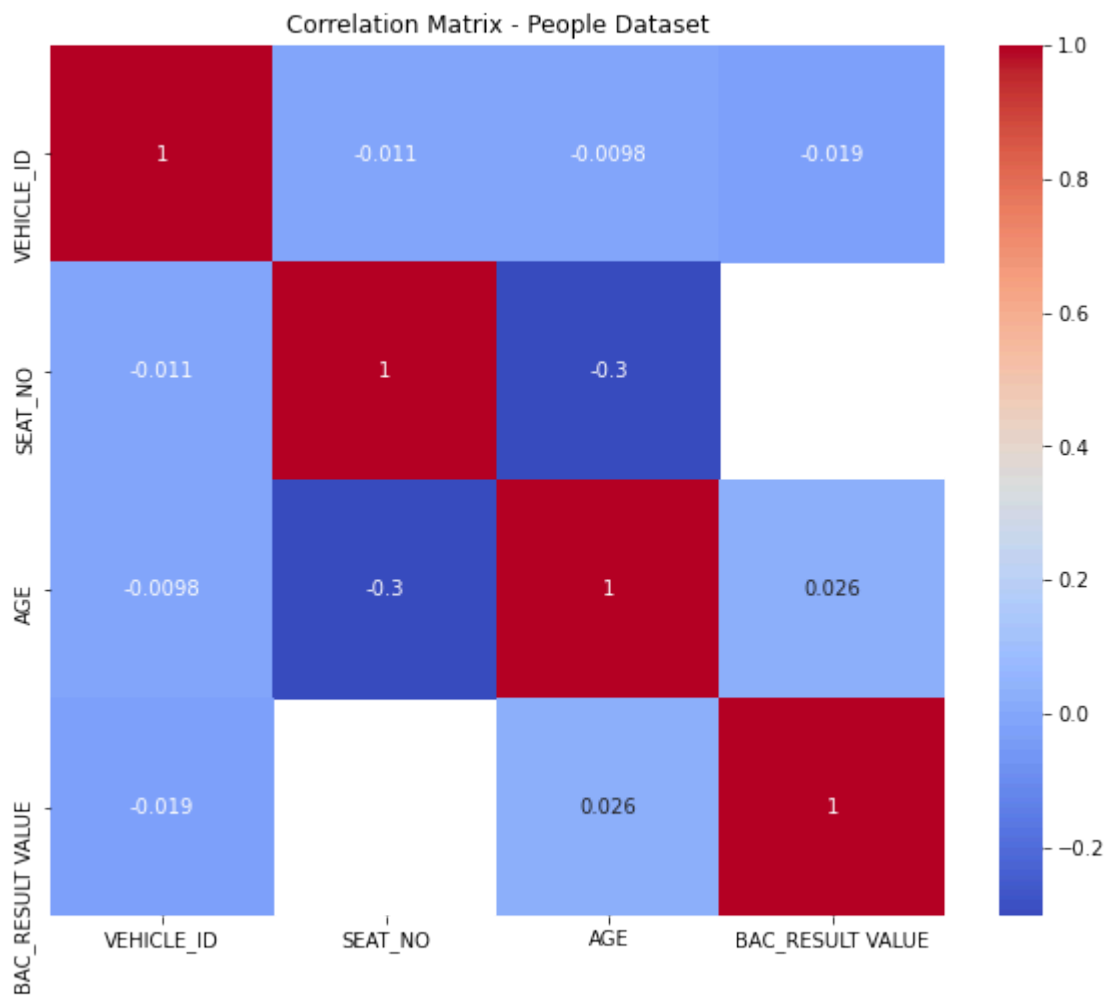
```
In [19]: # Scatter plot to examine relationships between numerical features in Crash
plt.figure(figsize=(8, 6))
sns.scatterplot(x='POSTED_SPEED_LIMIT', y='NUM_UNITS', data=crashes)
plt.title('Scatter Plot - Speed Limit vs Number of Units Involved (Crashes)')
plt.show()
```



```
In [20]: # Correlation matrix for numerical features in Crashes Dataset
plt.figure(figsize=(10, 8))
corr_matrix_crashes = crashes.corr()
sns.heatmap(corr_matrix_crashes, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix - Crashes Dataset')
plt.show()
```



```
In [21]: # Correlation matrix for numerical features in People Dataset
plt.figure(figsize=(10, 8))
corr_matrix_people = people.corr()
sns.heatmap(corr_matrix_people, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix - People Dataset')
plt.show()
```




```
In [22]: # Cross-tabulation for categorical features in People Dataset
cross_tab_people = pd.crosstab(people['PERSON_TYPE'], people['INJURY_CLASSIFICATION'])
print("\nCross-tabulation of 'PERSON_TYPE' and 'INJURY_CLASSIFICATION' in People Dataset")
```

Cross-tabulation of 'PERSON_TYPE' and 'INJURY_CLASSIFICATION' in People Dataset:

INJURY_CLASSIFICATION \ PERSON_TYPE	FATAL	INCAPACITATING INJURY	NO INDICATION OF INJURY
BICYCLE	31	943	3
DRIVER	405	6595	1072
NON-CONTACT VEHICLE	0	0	0
NON-MOTOR VEHICLE	3	22	0
PASSENGER	144	2960	263
PEDESTRIAN	206	2839	2

INJURY_CLASSIFICATION \ PERSON_TYPE	NONINCAPACITATING INJURY	REPORTED, NOT EVIDENT
BICYCLE	4926	1140
DRIVER	37093	23391
NON-CONTACT VEHICLE	1	0
NON-MOTOR VEHICLE	58	21
PASSENGER	19140	13010
PEDESTRIAN	8781	2539


```
In [23]: #Correlation matrix for numerical features in Crashes Dataset
plt.figure(figsize=(10, 8))
corr_matrix_vehicles = vehicles.corr()
sns.heatmap(corr_matrix_vehicles, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix - Vehicles Dataset')
plt.show()
```



```
In [24]: # Cross-tabulation for categorical features in Vehicles Dataset
cross_tab_vehicles = pd.crosstab(vehicles['VEHICLE_TYPE'], crashes['CRASH_TYPE'])
print("\nCross-tabulation of 'VEHICLE_TYPE' and 'CRASH_TYPE' in Vehicles Dataset")
```

Cross-tabulation of 'VEHICLE_TYPE' and 'CRASH_TYPE' in Vehicles Dataset:

VEHICLE_TYPE \ CRASH_TYPE	NO INJURY / DRIVE AWAY	INJURY AND / OR TOW DUE TO CRASH
---------------------------	------------------------	----------------------------------

3-WHEELED MOTORCYCLE (2 REAR WHEELS)	11
ALL-TERRAIN VEHICLE (ATV)	26
AUTOCYCLE	74
BUS OVER 15 PASS.	2222
BUS UP TO 15 PASS.	775
FARM EQUIPMENT	11
MOPED OR MOTORIZED BICYCLE	102
MOTOR DRIVEN CYCLE	30
MOTORCYCLE (OVER 150CC)	533
OTHER	2843
OTHER VEHICLE WITH TRAILER	335
PASSENGER	142214
PICKUP	7746
RECREATIONAL OFF-HIGHWAY VEHICLE (ROV)	3
SINGLE UNIT TRUCK WITH TRAILER	459
SNOWMOBILE	1
SPORT UTILITY VEHICLE (SUV)	32520
TRACTOR W/ SEMI-TRAILER	2065
TRACTOR W/O SEMI-TRAILER	259
TRUCK - SINGLE UNIT	4009
UNKNOWN/NA	21034
VAN/MINI-VAN	10165

VEHICLE_TYPE \ CRASH_TYPE	NO INJURY / DRIVE AWAY
3-WHEELED MOTORCYCLE (2 REAR WHEELS)	31
ALL-TERRAIN VEHICLE (ATV)	74
AUTOCYCLE	185
BUS OVER 15 PASS.	6097
BUS UP TO 15 PASS.	2092
FARM EQUIPMENT	23
MOPED OR MOTORIZED BICYCLE	281
MOTOR DRIVEN CYCLE	84
MOTORCYCLE (OVER 150CC)	1499
OTHER	7941
OTHER VEHICLE WITH TRAILER	788
PASSENGER	387267
PICKUP	21073
RECREATIONAL OFF-HIGHWAY VEHICLE (ROV)	15
SINGLE UNIT TRUCK WITH TRAILER	1190
SNOWMOBILE	2
SPORT UTILITY VEHICLE (SUV)	89083
TRACTOR W/ SEMI-TRAILER	5532
TRACTOR W/O SEMI-TRAILER	708
TRUCK - SINGLE UNIT	11390
UNKNOWN/NA	57099
VAN/MINI-VAN	27633

Missing Values

Objective: Identify missing data and assess the percentage of missing values per feature.

Identifying Missing Values

```
In [25]: # Identify missing values
print(crashes.isnull().sum())
```

CRASH_RECORD_ID	0
CRASH_DATE_EST_I	802055
CRASH_DATE	0
POSTED_SPEED_LIMIT	0
TRAFFIC_CONTROL_DEVICE	0
DEVICE_CONDITION	0
WEATHER_CONDITION	0
LIGHTING_CONDITION	0
FIRST_CRASH_TYPE	0
TRAFFICWAY_TYPE	0
LANE_CNT	667396
ALIGNMENT	0
ROADWAY_SURFACE_COND	0
ROAD_DEFECT	0
REPORT_TYPE	26403
CRASH_TYPE	0
INTERSECTION_RELATED_I	667808
NOT_RIGHT_OF_WAY_I	826744
HIT_AND_RUN_I	594832
DAMAGE	0
DATE_POLICE_NOTIFIED	0
PRIM_CONTRIBUTORY_CAUSE	0
SEC_CONTRIBUTORY_CAUSE	0
STREET_NO	0
STREET_DIRECTION	4
STREET_NAME	1
BEAT_OF_OCCURRENCE	5
PHOTOS_TAKEN_I	854736
STATEMENTS_TAKEN_I	846649
DOORING_I	863692
WORK_ZONE_I	861513
WORK_ZONE_TYPE	862631
WORKERS_PRESENT_I	865156
NUM_UNITS	0
MOST_SEVERE_INJURY	1916
INJURIES_TOTAL	1903
INJURIES_FATAL	1903
INJURIES_INCAPACITATING	1903
INJURIES_NON_INCAPACITATING	1903
INJURIES_REPORTED_NOT_EVIDENT	1903
INJURIES_NO_INDICATION	1903
INJURIES_UNKNOWN	1903
CRASH_HOUR	0
CRASH_DAY_OF_WEEK	0
CRASH_MONTH	0
LATITUDE	6138
LONGITUDE	6138
LOCATION	6138

dtype: int64

```
In [26]: print(people.isnull().sum())
```

```
PERSON_ID          0
PERSON_TYPE        4
CRASH_RECORD_ID    4
VEHICLE_ID        28804
CRASH_DATE         5
SEAT_NO          1167917
CITY              395404
STATE             381164
ZIPCODE           487700
SEX               22984
AGE              426196
DRIVERS_LICENSE_STATE 605557
DRIVERS_LICENSE_CLASS 736377
SAFETY_EQUIPMENT   4149
AIRBAG_DEPLOYED   27652
EJECTION          17696
INJURY_CLASSIFICATION 644
HOSPITAL          1212552
EMS_AGENCY        1308279
EMS_RUN_NO        1440923
DRIVER_ACTION     301796
DRIVER_VISION     302212
PHYSICAL_CONDITION 300982
PEDPEDAL_ACTION   1439792
PEDPEDAL_VISIBILITY 1439851
PEDPEDAL_LOCATION 1439797
BAC_RESULT        300779
BAC_RESULT VALUE  1465274
CELL_PHONE_USE    1465892
dtype: int64
```

```
In [27]: print(vehicles.isnull().sum())
```

```
CRASH_UNIT_ID      0
CRASH_RECORD_ID    0
CRASH_DATE         0
UNIT_NO            0
UNIT_TYPE          2209
...
CARGO_BODY_TYPE    1750780
LOAD_TYPE          1751403
HAZMAT_OUT_OF_SERVICE_I 1752567
MCS_OUT_OF_SERVICE_I 1752327
HAZMAT_CLASS       1763763
Length: 71, dtype: int64
```

Calculate the percentage of missing values

```
In [28]: crashes_missing_percentage = (crashes.isnull().sum() / len(crashes)) * 100
print("Crashes Missing Data Percentage:\n", crashes_missing_percentage)
```

Crashes Missing Data Percentage:

CRASH_RECORD_ID	0.000000
CRASH_DATE_EST_I	92.572116
CRASH_DATE	0.000000
POSTED_SPEED_LIMIT	0.000000
TRAFFIC_CONTROL_DEVICE	0.000000
DEVICE_CONDITION	0.000000
WEATHER_CONDITION	0.000000
LIGHTING_CONDITION	0.000000
FIRST_CRASH_TYPE	0.000000
TRAFFICWAY_TYPE	0.000000
LANE_CNT	77.029955
ALIGNMENT	0.000000
ROADWAY_SURFACE_COND	0.000000
ROAD_DEFECT	0.000000
REPORT_TYPE	3.047399
CRASH_TYPE	0.000000
INTERSECTION_RELATED_I	77.077507
NOT_RIGHT_OF_WAY_I	95.421688
HIT_AND_RUN_I	68.654715
DAMAGE	0.000000
DATE_POLICE_NOTIFIED	0.000000
PRIM_CONTRIBUTORY_CAUSE	0.000000
SEC_CONTRIBUTORY_CAUSE	0.000000
STREET_NO	0.000000
STREET_DIRECTION	0.000462
STREET_NAME	0.000115
BEAT_OF_OCCURRENCE	0.000577
PHOTOS_TAKEN_I	98.652487
STATEMENTS_TAKEN_I	97.719096
DOORING_I	99.686177
WORK_ZONE_I	99.434679
WORK_ZONE_TYPE	99.563717
WORKERS_PRESENT_I	99.855150
NUM_UNITS	0.000000
MOST_SEVERE_INJURY	0.221142
INJURIES_TOTAL	0.219642
INJURIES_FATAL	0.219642
INJURIES_INCAPACITATING	0.219642
INJURIES_NON_INCAPACITATING	0.219642
INJURIES_REPORTED_NOT_EVIDENT	0.219642
INJURIES_NO_INDICATION	0.219642
INJURIES_UNKNOWN	0.219642
CRASH_HOUR	0.000000
CRASH_DAY_OF_WEEK	0.000000
CRASH_MONTH	0.000000
LATITUDE	0.708440
LONGITUDE	0.708440
LOCATION	0.708440

dtype: float64

```
In [29]: people_missing_percentage = (people.isnull().sum() / len(people)) * 100
print("People Missing Data Percentage:\n", people_missing_percentage)
```

```
People Missing Data Percentage:
PERSON_ID          0.000000
PERSON_TYPE        0.000273
CRASH_RECORD_ID    0.000273
VEHICLE_ID         1.963397
CRASH_DATE         0.000341
SEAT_NO           79.609952
CITY              26.952338
STATE            25.981682
ZIPCODE          33.243607
SEX              1.566683
AGE             29.051245
DRIVERS_LICENSE_STATE 41.277217
DRIVERS_LICENSE_CLASS 50.194438
SAFETY_EQUIPMENT    0.282813
AIRBAG_DEPLOYED    1.884872
EJECTION           1.206231
INJURY_CLASSIFICATION 0.043898
HOSPITAL          82.652454
EMS_AGENCY        89.177594
EMS_RUN_NO        98.219146
DRIVER_ACTION     20.571637
DRIVER_VISION     20.599994
PHYSICAL_CONDITION 20.516152
PEDPEDAL_ACTION   98.142053
PEDPEDAL_VISIBILITY 98.146074
PEDPEDAL_LOCATION 98.142393
BAC_RESULT        20.502315
BAC_RESULT VALUE   99.879009
CELL_PHONE_USE    99.921134
dtype: float64
```

```
In [30]: vehicles_missing_percentage = (vehicles.isnull().sum() / len(vehicles)) * 100
print("Vehicles Missing Data Percentage:\n", vehicles_missing_percentage)
```

```
Vehicles Missing Data Percentage:
CRASH_UNIT_ID      0.000000
CRASH_RECORD_ID    0.000000
CRASH_DATE         0.000000
UNIT_NO           0.000000
UNIT_TYPE          0.125163
...
CARGO_BODY_TYPE    99.199955
LOAD_TYPE          99.235254
HAZMAT_OUT_OF_SERVICE_I 99.301207
MCS_OUT_OF_SERVICE_I 99.287608
HAZMAT_CLASS       99.935577
Length: 71, dtype: float64
```

Dropping Columns with High % of Missing Values

Columns to Keep Despite Missing Values

People Dataset

DRIVERS_LICENSE_CLASS (50.19% missing): Could be important for understanding driver qualifications.

Crashes Dataset

REPORT_TYPE (3.05% missing): Low enough missing values that it might be worth keeping.

```
In [31]: # Crashes dataset
crashes_cleaned = crashes.drop(columns=['CRASH_DATE_EST_I', 'LANE_CNT', 'IN
```

```
In [32]: # People dataset
people_cleaned = people.drop(columns=['SEAT_NO', 'HOSPITAL', 'EMS_AGENCY',
```

```
In [33]: # Vehicles dataset
vehicles_cleaned = vehicles.drop(columns=['CARGO_BODY_TYPE', 'LOAD_TYPE', '
```

Save the cleaned datasets to the processed_data folder

```
In [34]: #Save the cleaned datasets to the processed_data folder

output_folder = 'C:/Users/MNJORGE16/Desktop/Moringa/phase_3/project__phase

crashes_cleaned.to_csv(output_folder + 'cleaned_crashes.csv', index=False)
people_cleaned.to_csv(output_folder + 'cleaned_people.csv', index=False)
vehicles_cleaned.to_csv(output_folder + 'cleaned_vehicles.csv', index=False)

print("Cleaned datasets saved successfully to the processed_data folder.")
```

Cleaned datasets saved successfully to the processed_data folder.

Summary Statistics After Dropping Columns

Summary statistics for numerical and categorical variables to understand their distributions.


```
In [35]: # Summary statistics for numerical features - Crashes  
print("Crashes Summary Statistics After Dropping Missing Values:\n", crashe
```

Crashes Summary Statistics After Dropping Missing Values:

	POSTED_SPEED_LIMIT	STREET_NO	BEAT_OF_OCCURRENCE	NUM_UN
ITS \				
count	866411.000000	866411.000000	866406.000000	866411.0000
00				
mean	28.415733	3687.152034	1244.469227	2.0351
17				
std	6.131785	2882.599171	705.126059	0.4527
53				
min	0.000000	0.000000	111.000000	1.0000
00				
25%	30.000000	1250.000000	714.000000	2.0000
00				
50%	30.000000	3201.000000	1212.000000	2.0000
00				
75%	30.000000	5580.000000	1822.000000	2.0000
00				
max	99.000000	451100.000000	6100.000000	18.0000
00				

	INJURIES_TOTAL	INJURIES_FATAL	INJURIES_INCAPACITATING \
count	864508.000000	864508.000000	864508.000000
mean	0.192690	0.001194	0.019823
std	0.570222	0.037455	0.164843
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	21.000000	4.000000	10.000000

	INJURIES_NON_INCAPACITATING	INJURIES_REPORTED_NOT_EVIDENT \
count	864508.000000	864508.000000
mean	0.108248	0.063426
std	0.424294	0.323900
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	21.000000	15.000000

	INJURIES_NO_INDICATION	INJURIES_UNKNOWN	CRASH_HOUR \
count	864508.000000	864508.0	866411.000000
mean	2.001795	0.0	13.205135
std	1.157261	0.0	5.573549
min	0.000000	0.0	0.000000
25%	1.000000	0.0	9.000000
50%	2.000000	0.0	14.000000
75%	2.000000	0.0	17.000000
max	61.000000	0.0	23.000000

	CRASH_DAY_OF_WEEK	CRASH_MONTH	LATITUDE	LONGITUDE
count	866411.000000	866411.000000	860273.000000	860273.000000
mean	4.122962	6.606381	41.855078	-87.673657
std	1.981495	3.377482	0.333591	0.677645
min	1.000000	1.000000	0.000000	-87.936193
25%	2.000000	4.000000	41.782879	-87.721774
50%	4.000000	7.000000	41.874945	-87.674177
75%	6.000000	10.000000	41.924490	-87.633463
max	7.000000	12.000000	42.022780	0.000000

```
In [36]: # Summary statistics for numerical features - People
print("People Summary Statistics After Dropping Missing Values:\n", people_
```

People Summary Statistics After Dropping Missing Values:

	VEHICLE_ID	AGE
count	1.438245e+06	1.040853e+06
mean	6.905905e+05	3.781694e+01
std	4.038528e+05	1.710846e+01
min	2.000000e+00	-1.770000e+02
25%	3.444260e+05	2.500000e+01
50%	6.826620e+05	3.500000e+01
75%	1.034161e+06	5.000000e+01
max	1.801497e+06	1.100000e+02

```
In [37]: # Summary statistics for numerical features - Vehicles
print("Vehicles Summary Statistics After Dropping Missing Values:\n", vehic
```

Vehicles Summary Statistics After Dropping Missing Values:

	CRASH_UNIT_ID	UNIT_NO	NUM_PASSENGERS	VEHICLE_ID \
count	1.764900e+06	1.764900e+06	261313.000000	1.724034e+06
mean	9.438774e+05	3.705683e+00	1.470750	8.976615e+05
std	5.463825e+05	2.843844e+03	1.055718	5.186095e+05
min	2.000000e+00	0.000000e+00	1.000000	2.000000e+00
25%	4.698658e+05	1.000000e+00	1.000000	4.489182e+05
50%	9.450715e+05	2.000000e+00	1.000000	8.959025e+05
75%	1.417380e+06	2.000000e+00	2.000000	1.346214e+06
max	1.888827e+06	3.778035e+06	59.000000	1.799377e+06

	VEHICLE_YEAR	OCCUPANT_CNT	CMV_ID	TRAILER1_LENGTH \
count	1.448829e+06	1.724034e+06	17859.000000	2393.000000
mean	2.014207e+03	1.079142e+00	9960.036564	48.511910
std	1.385204e+02	7.815274e-01	5757.641443	20.695514
min	1.900000e+03	0.000000e+00	1.000000	1.000000
25%	2.007000e+03	1.000000e+00	4918.500000	45.000000
50%	2.013000e+03	1.000000e+00	9988.000000	53.000000
75%	2.017000e+03	1.000000e+00	14967.500000	53.000000
max	9.999000e+03	9.900000e+01	19878.000000	740.000000

	TRAILER2_LENGTH	TOTAL_VEHICLE_LENGTH	AXLE_CNT
count	70.000000	2918.000000	4396.000000
mean	44.271429	53.225497	9.619882
std	28.008240	31.291466	392.233256
min	1.000000	1.000000	1.000000
25%	24.250000	35.000000	2.000000
50%	50.000000	53.000000	3.000000
75%	53.000000	66.000000	5.000000
max	123.000000	999.000000	26009.000000

Merging the three data sets

Merging the crashes_cleaned, people_cleaned, and vehicles_cleaned datasets on the common key (CRASH_RECORD_ID).

```
In [38]: #Merge the cleaned datasets on 'CRASH_RECORD_ID'
merged_df = pd.merge(pd.merge(crashes_cleaned, people_cleaned, on='CRASH_RECORD_ID'),
                      #Save the merged dataset to the specified folder
                      output_path = 'C:/Users/MNJOROG16/Desktop/Moringa/phase_3/project__phase3/Project-Ph3-Chicago-Car-Crashes-Predictions/data/processed_data/cleaned_merged_traffic_crashes.csv',
                      merged_df.to_csv(output_path, index=False))

print(f"Cleaned merged dataset saved successfully to {output_path}")
```

Cleaned merged dataset saved successfully to C:/Users/MNJOROG16/Desktop/Moringa/phase_3/project__phase3/Project-Ph3-Chicago-Car-Crashes-Predictions/data/processed_data/cleaned_merged_traffic_crashes.csv

Inspection of the Merged Dataset

```
In [39]: # Inspect the first few rows of the merged dataset
print(merged_df.head())
```

		CRASH_RECORD_ID	CRASH_DAT
E_x \			
0	004cd14d0303a9163aad69a2d7f341b7da2a8572b2ab33...	11/26/2019	08:38:00
AM			
1	004cd14d0303a9163aad69a2d7f341b7da2a8572b2ab33...	11/26/2019	08:38:00
AM			
2	004cd14d0303a9163aad69a2d7f341b7da2a8572b2ab33...	11/26/2019	08:38:00
AM			
3	004cd14d0303a9163aad69a2d7f341b7da2a8572b2ab33...	11/26/2019	08:38:00
AM			
4	359bf9f5872d646bb63576e55b1e0b480dc93c2b935ab5...	01/31/2022	07:45:00
PM			

	POSTED_SPEED_LIMIT	TRAFFIC_CONTROL_DEVICE	DEVICE_CONDITION	\
0	25	NO CONTROLS	NO CONTROLS	
1	25	NO CONTROLS	NO CONTROLS	
2	25	NO CONTROLS	NO CONTROLS	
3	25	NO CONTROLS	NO CONTROLS	
4	25	NO CONTROLS	NO CONTROLS	

	WEATHER_CONDITION	LIGHTING_CONDITION	FIRST_CRASH_TYPE	TRAFFICWAY_TYPE
\				
0	CLEAR	DAYLIGHT	PEDESTRIAN	ONE-WAY
1	CLEAR	DAYLIGHT	PEDESTRIAN	ONE-WAY
2	CLEAR	DAYLIGHT	PEDESTRIAN	ONE-WAY
3	CLEAR	DAYLIGHT	PEDESTRIAN	ONE-WAY
4	CLEAR	DARKNESS	REAR END	ONE-WAY

	ALIGNMENT	...	MCS_VIO_CAUSE_CRASH_I	IDOT_PERMIT_NO	WIDE_LOAD
_I \					
0	CURVE ON GRADE	...	NaN	NaN	N
aN					
1	CURVE ON GRADE	...	NaN	NaN	N
aN					
2	CURVE ON GRADE	...	NaN	NaN	N
aN					
3	CURVE ON GRADE	...	NaN	NaN	N
aN					
4	STRAIGHT AND LEVEL	...	NaN	NaN	N
aN					

	TRAILER1_WIDTH	TRAILER2_WIDTH	TRAILER1_LENGTH	TRAILER2_LENGTH	\
0	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	

	TOTAL_VEHICLE_LENGTH	AXLE_CNT	VEHICLE_CONFIG
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN

[5 rows x 129 columns]

```
In [40]: # Verify the number of rows and columns
print("Merged Dataset Shape:", merged_df.shape)
```

Merged Dataset Shape: (3076376, 129)

```
In [41]: # Check the column names and data types
print("Merged Dataset Columns and Data Types:\n", merged_df.dtypes)
```

Merged Dataset Columns and Data Types:

CRASH_RECORD_ID	object
CRASH_DATE_x	object
POSTED_SPEED_LIMIT	int64
TRAFFIC_CONTROL_DEVICE	object
DEVICE_CONDITION	object
...	
TRAILER1_LENGTH	float64
TRAILER2_LENGTH	float64
TOTAL_VEHICLE_LENGTH	float64
AXLE_CNT	float64
VEHICLE_CONFIG	object

Length: 129, dtype: object

Handling Missing Values- Merged Dataset

Identifying and address any missing values that may affect the analysis and modeling process.

```
In [42]: # Identify missing values in the merged dataset
missing_values = merged_df.isnull().sum()

print (missing_values)
```

CRASH_RECORD_ID	0
CRASH_DATE_x	0
POSTED_SPEED_LIMIT	0
TRAFFIC_CONTROL_DEVICE	0
DEVICE_CONDITION	0
...	
TRAILER1_LENGTH	3072274
TRAILER2_LENGTH	3076240
TOTAL_VEHICLE_LENGTH	3070784
AXLE_CNT	3068380
VEHICLE_CONFIG	3050604

Length: 129, dtype: int64

```
In [43]: # Calculate the percentage of missing values for each feature
missing_percentage = (missing_values / len(merged_df)) * 100

# Display features with missing values
print("Missing Values in Merged Dataset:\n", missing_percentage[missing_per
```

```
Missing Values in Merged Dataset:
REPORT_TYPE          3.832431
STREET_DIRECTION     0.000455
STREET_NAME          0.000130
BEAT_OF_OCCURRENCE   0.000780
PHOTOS_TAKEN_I       98.674317
...
TRAILER1_LENGTH      99.866661
TRAILER2_LENGTH      99.995579
TOTAL_VEHICLE_LENGTH 99.818228
AXLE_CNT             99.740084
VEHICLE_CONFIG       99.162261
Length: 94, dtype: float64
```

Drop Features with Extremely High Missing Values

Criteria: Features with more than 90% missing data are typically considered for removal unless they are critical.

Action: Drop features like PHOTOS_TAKEN_I, TRAILER1_LENGTH, TRAILER2_LENGTH, TOTAL_VEHICLE_LENGTH, AXLE_CNT, and VEHICLE_CONFIG because their missing rates are extremely high (> 99%).

```
In [44]: # Drop features with more than 90% missing values
columns_to_drop = ['PHOTOS_TAKEN_I', 'TRAILER1_LENGTH', 'TRAILER2_LENGTH',
merged_df_dropped = merged_df.drop(columns=columns_to_drop)

print(f"Dropped columns: {columns_to_drop}")
print("New dataset shape after dropping columns:", merged_df_dropped.shape)
```

```
Dropped columns: ['PHOTOS_TAKEN_I', 'TRAILER1_LENGTH', 'TRAILER2_LENGTH',
'TOTAL_VEHICLE_LENGTH', 'AXLE_CNT', 'VEHICLE_CONFIG']
New dataset shape after dropping columns: (3076376, 123)
```

Impute Missing Values for Important Features

Criteria: For features with lower missing rates (e.g., REPORT_TYPE, STREET_DIRECTION, STREET_NAME, BEAT_OF_OCCURRENCE), imputation is appropriate.

Imputation Methods: Categorical Features: Impute missing values using the mode (most frequent value).

Numerical Features: If any were present, you could use mean, median, or other statistical methods.

```
In [45]: # Impute missing values for categorical features with the mode
merged_df_dropped['REPORT_TYPE'].fillna(merged_df_dropped['REPORT_TYPE'].mode[0])
merged_df_dropped['STREET_DIRECTION'].fillna(merged_df_dropped['STREET_DIRECTION'].mode[0])
merged_df_dropped['STREET_NAME'].fillna(merged_df_dropped['STREET_NAME'].mode[0])
merged_df_dropped['BEAT_OF_OCCURRENCE'].fillna(merged_df_dropped['BEAT_OF_OCCURRENCE'].mode[0])

# Check if there are any remaining missing values after imputation
print("Remaining Missing Values After Imputation:\n", merged_df_dropped.isnull().sum())
```

Remaining Missing Values After Imputation:

STATEMENTS_TAKEN_I	2998531
DOORING_I	3066855
WORK_ZONE_I	3058959
WORK_ZONE_TYPE	3062786
WORKERS_PRESENT_I	3072073
	...
MCS_VIO_CAUSE_CRASH_I	3054541
IDOT_PERMIT_NO	3074778
WIDE_LOAD_I	3076133
TRAILER1_WIDTH	3071381
TRAILER2_WIDTH	3075765

Length: 84, dtype: int64

```
In [46]: # Identify columns with a high percentage of missing values (e.g., >90%)
high_missing_columns = merged_df_dropped.columns[merged_df_dropped.isnull().sum() > 0.9 * merged_df_dropped.shape[0]]

print (high_missing_columns)
```

```
Index(['STATEMENTS_TAKEN_I', 'DOORING_I', 'WORK_ZONE_I', 'WORK_ZONE_TYP
E',
      'WORKERS_PRESENT_I', 'PEDPEDAL_VISIBILITY', 'PEDPEDAL_LOCATION',
      'CMRC_VEH_I', 'FIRE_I', 'EXCEED_SPEED_LIMIT_I', 'TOWED_TO', 'AREA_
00_I',
      'AREA_03_I', 'AREA_04_I', 'AREA_09_I', 'AREA_10_I', 'AREA_99_I',
      'CMV_ID', 'USDOT_NO', 'CCMC_NO', 'ILCC_NO', 'COMMERCIAL_SRC', 'GVW
R',
      'CARRIER_NAME', 'CARRIER_STATE', 'CARRIER_CITY', 'HAZMAT_PLACARDS_
I',
      'HAZMAT_NAME', 'UN_NO', 'HAZMAT_PRESENT_I', 'HAZMAT_REPORT_I',
      'HAZMAT_REPORT_NO', 'MCS_REPORT_I', 'MCS_REPORT_NO',
      'HAZMAT_VIO_CAUSE_CRASH_I', 'MCS_VIO_CAUSE_CRASH_I', 'IDOT_PERMIT_
NO',
      'WIDE_LOAD_I', 'TRAILER1_WIDTH', 'TRAILER2_WIDTH'],
      dtype='object')
```



```
In [47]: # Drop these columns
merged_df_final = merged_df_dropped.drop(columns=high_missing_columns)

print(f"Dropped columns with high missing values: {high_missing_columns}")
print("New dataset shape after dropping high-missing-value columns:", merged_df_final.shape)
```

```
Dropped columns with high missing values: Index(['STATEMENTS_TAKEN_I', 'D
OORING_I', 'WORK_ZONE_I', 'WORK_ZONE_TYPE',
        'WORKERS_PRESENT_I', 'PEDPEDAL_VISIBILITY', 'PEDPEDAL_LOCATION',
        'CMRC_VEH_I', 'FIRE_I', 'EXCEED_SPEED_LIMIT_I', 'TOWED_TO', 'AREA_
00_I',
        'AREA_03_I', 'AREA_04_I', 'AREA_09_I', 'AREA_10_I', 'AREA_99_I',
        'CMV_ID', 'USDOT_NO', 'CCMC_NO', 'ILCC_NO', 'COMMERCIAL_SRC', 'GVW
R',
        'CARRIER_NAME', 'CARRIER_STATE', 'CARRIER_CITY', 'HAZMAT_PLACARDS_
I',
        'HAZMAT_NAME', 'UN_NO', 'HAZMAT_PRESENT_I', 'HAZMAT_REPORT_I',
        'HAZMAT_REPORT_NO', 'MCS_REPORT_I', 'MCS_REPORT_NO',
        'HAZMAT_VIO_CAUSE_CRASH_I', 'MCS_VIO_CAUSE_CRASH_I', 'IDOT_PERMIT_
NO',
        'WIDE_LOAD_I', 'TRAILER1_WIDTH', 'TRAILER2_WIDTH'],
        dtype='object')
New dataset shape after dropping high-missing-value columns: (3076376, 83)
```

```
In [48]: # Check for any remaining missing values in the final dataset
remaining_missing_values = merged_df_final.isnull().sum()
remaining_missing_percentage = (remaining_missing_values / len(merged_df_final)) * 100

# Display features with remaining missing values
print("Remaining Missing Values After Dropping High-Missing-Value Columns:\n")
```

Remaining Missing Values After Dropping High-Missing-Value Columns:

MOST_SEVERE_INJURY	0.000683
LATITUDE	0.619723
LONGITUDE	0.619723
LOCATION	0.619723
VEHICLE_ID_x	2.039608
CITY	26.815545
STATE	25.820706
ZIPCODE	33.082497
SEX	1.587745
AGE	28.843938
DRIVERS_LICENSE_STATE	41.419287
DRIVERS_LICENSE_CLASS	50.295575
SAFETY_EQUIPMENT	0.309195
AIRBAG_DEPLOYED	1.936369
EJECTION	1.269481
INJURY_CLASSIFICATION	0.050709
DRIVER_ACTION	20.782083
DRIVER_VISION	20.812768
PHYSICAL_CONDITION	20.722955
BAC_RESULT	20.705824
UNIT_TYPE	0.119264
NUM_PASSENGERS	74.837504
VEHICLE_ID_y	2.310738
MAKE	2.311128
MODEL	2.323578
LIC_PLATE_STATE	9.615665
VEHICLE_YEAR	15.715277
VEHICLE_DEFECT	2.310738
VEHICLE_TYPE	2.310738
VEHICLE_USE	2.310738
TRAVEL_DIRECTION	2.310738
MANEUVER	2.310738
TOWED_I	86.295888
OCCUPANT_CNT	2.310738
TOWED_BY	89.707305
AREA_01_I	71.485215
AREA_02_I	83.198900
AREA_05_I	84.180347
AREA_06_I	84.124892
AREA_07_I	86.402767
AREA_08_I	84.885593
AREA_11_I	82.518782
AREA_12_I	82.184297
FIRST_CONTACT_POINT	2.517768

dtype: float64

```
In [49]: # Impute categorical features with the mode
for column in ['MOST_SEVERE_INJURY', 'SEX', 'DRIVERS_LICENSE_STATE', 'DRIVE',
               'AIRBAG_DEPLOYED', 'EJECTION', 'INJURY_CLASSIFICATION', 'DRI',
               'PHYSICAL_CONDITION', 'UNIT_TYPE', 'MAKE', 'FIRST_CONTACT_PC',
               'PHYSICAL_CONDITION', 'UNIT_TYPE', 'MAKE', 'FIRST_CONTACT_PC']:
    merged_df_final[column].fillna(merged_df_final[column].mode()[0], inplace=True)

# Impute numerical features with the median
for column in ['LATITUDE', 'LONGITUDE', 'AGE']:
    merged_df_final[column].fillna(merged_df_final[column].median(), inplace=True)
```

```
In [50]: # Check for any remaining missing values
remaining_missing_values = merged_df_final.isnull().sum()
print("Remaining missing values:\n", remaining_missing_values[remaining_mis

print (remaining_missing_values)
```

```
Remaining missing values:
  LOCATION      19065
VEHICLE_ID_x    62746
CITY           824947
STATE          794342
ZIPCODE        1017742
BAC_RESULT      636989
NUM_PASSENGERS  2302283
VEHICLE_ID_y    71087
MODEL          71482
LIC_PLATE_STATE 295814
VEHICLE_YEAR    483461
VEHICLE_DEFECT  71087
VEHICLE_TYPE    71087
VEHICLE_USE     71087
TRAVEL_DIRECTION 71087
MANEUVER        71087
TOWED_I         2654786
OCCUPANT_CNT    71087
TOWED_BY        2759734
AREA_01_I       2199154
AREA_02_I       2559511
AREA_05_I       2589704
AREA_06_I       2587998
AREA_07_I       2658074
AREA_08_I       2611400
AREA_11_I       2538588
AREA_12_I       2528298
dtype: int64
CRASH_RECORD_ID      0
CRASH_DATE_x         0
POSTED_SPEED_LIMIT   0
TRAFFIC_CONTROL_DEVICE 0
DEVICE_CONDITION     0

...
AREA_07_I      2658074
AREA_08_I      2611400
AREA_11_I      2538588
AREA_12_I      2528298
FIRST_CONTACT_POINT 0
Length: 83, dtype: int64
```

Analyze Feature Importance and Missing Data Percentage

For each feature, consider its importance to the model and the percentage of missing data. This will guide whether to impute or drop the feature.

High Importance & Low Missing Data (<20%): Impute missing values.

High Importance & High Missing Data (>20%): Consider imputation if the feature is crucial, otherwise consider dropping.

Low Importance & High Missing Data (>20%): Likely candidates for dropping.

Review of Missing Data Percentages

From the previous data, we have the following features with missing values and their percentages:

LOCATION: 0.62% missing

VEHICLE_ID_x: 2.04% missing

CITY: 26.82% missing

STATE: 25.82% missing

ZIPCODE: 33.08% missing

BAC_RESULT: 20.71% missing

NUM_PASSENGERS: 74.84% missing

VEHICLE_ID_y: 2.31% missing

MODEL: 2.32% missing

LIC_PLATE_STATE: 9.61% missing

VEHICLE_YEAR: 15.72% missing

VEHICLE_DEFECT: 2.31% missing

VEHICLE_TYPE: 2.31% missing

VEHICLE_USE: 2.31% missing

TRAVEL_DIRECTION: 2.31% missing

MANEUVER: 2.31% missing

TOWED_I: 86.34% missing

OCCUPANT_CNT: 2.31% missing

TOWED_BY: 89.70% missing

AREA_01_I - AREA_12_I: Varies, mostly >70% missing

High Importance Features based on their importance to (INJURY_CLASSIFICATION) and the percentage of missing data.

LOCATION (0.62%): High importance, low missing data. Impute.

VEHICLE_ID_x (2.04%) and VEHICLE_ID_y (2.31%): Moderate importance for identifying specific vehicles, low missing data. Impute.

MODEL (2.32%): High importance for determining vehicle type, low missing data. Impute.

LIC_PLATE_STATE (9.61%): High importance for location-based analysis, moderate missing data. Impute.

VEHICLE_YEAR (15.72%): High importance for determining vehicle age, moderate missing data. Impute.

VEHICLE_TYPE, VEHICLE_DEFECT, VEHICLE_USE, TRAVEL_DIRECTION, MANEUVER, OCCUPANT_CNT (All ~2.31%): High importance for understanding crash

Moderate to Low Importance Features

BAC_RESULT (20.71%): Important, but a higher percentage of missing data. Impute

CITY (26.82%) and STATE (25.82%): Important for location-based analysis but high missing data. Consider Dropping.

ZIPCODE (33.08%): Similar to CITY and STATE, consider dropping due to high missing data.

TOWED_I (86.34%): Low importance, very high missing data. Drop.

TOWED_BY (89.70%): Low importance, very high missing data. Drop.

AREA_01_I - AREA_12_I (>70%): Low importance, very high missing data. Drop

Further Imputing and Dropping

Impute

LOCATION

VEHICLE_ID_x, VEHICLE_ID_y

MODEL

LIC_PLATE_STATE

VEHICLE_YEAR

VEHICLE_TYPE, VEHICLE_DEFECT, VEHICLE_USE, TRAVEL_DIRECTION, MANEUVER, OCCUPANT_CNT

BAC_RESULT

Drop

CITY, STATE, ZIPCODE: Due to the high percentage of missing data, these features should be dropped unless you have strong reasons to keep them.

TOWED_I, TOWED_BY: Drop due to very high missing data and low importance.

AREA_01_I - AREA_12_I: Drop due to very high missing data and low importance.

```
In [51]: # Impute high importance features with missing data
for column in ['LOCATION', 'VEHICLE_ID_x', 'VEHICLE_ID_y', 'MODEL', 'LIC_PL',
               'VEHICLE_DEFECT', 'VEHICLE_TYPE', 'VEHICLE_USE', 'TRAVEL_DIST',
               'OCCUPANT_CNT', 'BAC_RESULT']:
    merged_df_final[column].fillna(merged_df_final[column].mode()[0], inplace=True)
```

```
In [52]: # Drop features with high missing data and low importance
columns_to_drop = ['CITY', 'STATE', 'ZIPCODE', 'TOWED_I', 'TOWED_BY',
                  'AREA_01_I', 'AREA_02_I', 'AREA_05_I', 'AREA_06_I',
                  'AREA_07_I', 'AREA_11_I', 'AREA_12_I']

merged_df_final = merged_df_final.drop(columns=columns_to_drop)

print(f"Dropped columns: {columns_to_drop}")
print("Dataset shape after dropping columns:", merged_df_final.shape)
```

Dropped columns: ['CITY', 'STATE', 'ZIPCODE', 'TOWED_I', 'TOWED_BY', 'AREA_01_I', 'AREA_02_I', 'AREA_05_I', 'AREA_06_I', 'AREA_07_I', 'AREA_11_I', 'AREA_12_I']
 Dataset shape after dropping columns: (3076376, 71)

```
In [53]: # Final check for remaining missing values
remaining_missing_values = merged_df_final.isnull().sum()
print("Remaining missing values:\n", remaining_missing_values[remaining_missing_values > 0])
```

Remaining missing values:
 NUM_PASSENGERS 2302283
 AREA_08_I 2611400
 dtype: int64

```
In [54]: # Drop the remaining features with high missing values
columns_to_drop = ['NUM_PASSENGERS', 'AREA_08_I']

merged_df_final = merged_df_final.drop(columns=columns_to_drop)
```

```
In [55]: # Final check for any remaining missing values
remaining_missing_values = merged_df_final.isnull().sum()
print("Final check for remaining missing values:\n", remaining_missing_values[remaining_missing_values > 0])
```

Final check for remaining missing values:
 Series([], dtype: int64)

```
In [56]: # Verify the shape of the final dataset
print("Final dataset shape after dropping remaining columns:", merged_df_final.shape)
```

Final dataset shape after dropping remaining columns: (3076376, 69)

Data Relationships


```
In [57]: # Define the numerical variables of interest
numerical_features = ['LATITUDE', 'LONGITUDE', 'AGE', 'POSTED_SPEED_LIMIT',
```

```
In [58]: # Define the path where the image will be saved
save_path = r'C:\Users\MNJOROGE16\Desktop\Moringa\phase_3\project__phase3\

# Correlation Matrix and Heatmap for Numerical Features
plt.figure(figsize=(12, 8))
corr_matrix = merged_df_final[numerical_features].corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix of Numerical Features')

# Save the figure as a PNG file
plt.savefig(save_path)

# Display the plot
plt.show()
```



Correlation Matrix Heatmap Interpretation

AGE Correlation with other features - AGE has almost no correlation with POSTED_SPEED_LIMIT (-0.02), NUM_UNITS (-0.03), LATITUDE (-0.00), and LONGITUDE (0.00).

Interpretation - Age does not appear to have a linear relationship with any of the other numerical features. This suggests that age might not directly influence or be influenced by these factors in the context of car crashes, making it a potentially independent variable in the model.

POSTED_SPEED_LIMIT Correlation with other features - POSTED_SPEED_LIMIT has very weak correlations with NUM_UNITS (0.06), LATITUDE (-0.01), and LONGITUDE (0.02).

Interpretation - The speed limit at the crash location shows minimal correlation with other numerical variables. This suggests that posted speed limits do not strongly interact with these variables, implying that speed limits might act independently in predicting injury severity.

NUM_UNITS Correlation with other features - NUM_UNITS has a weak correlation with POSTED_SPEED_LIMIT (0.06) and negligible correlation with other features.

Interpretation - The number of units involved in an accident (likely vehicles) is somewhat related to the speed limit, which makes sense, as different traffic conditions and regulations might influence both. However, the overall low correlations suggest that the number of units is generally independent of the other numerical features.

LATITUDE and LONGITUDE

Correlation with each other - There is a strong negative correlation between LATITUDE and LONGITUDE (-0.98).

Interpretation - The strong negative correlation between latitude and longitude indicates that these two variables are closely related, likely due to the geographical layout of the region covered in the dataset. This relationship could be important for geospatial analysis but might not directly impact injury classification unless location-based patterns are significant.

Conclusion

Low Correlation Among Most Features - Most of the numerical features have very low correlations with each other. This suggests that these variables operate relatively independently, which is useful for modeling as it reduces the risk of multicollinearity, which can distort the predictive power of individual features.

Latitude and Longitude - The strong correlation between LATITUDE and LONGITUDE highlights that these two variables are geographically dependent. This could be useful if you plan to include geospatial analysis or location-based features in your model. However, since their correlation is high, you might consider using one or combining them into a new feature to avoid redundancy.


```
In [59]: # Define the save path
save_path = r'C:\Users\MNJORGE16\Desktop\Moringa\phase_3\project__phase3\F

# Pairwise Scatter Plots for Numerical Features
pair_plot = sns.pairplot(merged_df_final[numerical_features])

# Adjust the title placement
pair_plot.fig.suptitle('Pairwise Scatter Plots of Numerical Features', y=1.

# Save the figure as a PNG file
pair_plot.savefig(save_path)

# Display the plot
plt.show()
```



Interpretation of Pairwise Scatter Plots for Numerical Features

AGE vs. Other Features

AGE vs. POSTED_SPEED_LIMIT

Observation - There is no clear linear relationship between AGE and POSTED_SPEED_LIMIT. The data points are scattered widely across different speed limits, regardless of age.

Interpretation - Age does not seem to influence the speed limit at which accidents occur, indicating these features are likely independent.

AGE vs. NUM_UNITS

Observation - The scatter plot shows a broad distribution with no distinct pattern. Most accidents involve 2-3 units across all ages.

Interpretation - The number of vehicles involved in an accident does not appear to be directly related to the age of individuals involved.

AGE vs. LATITUDE/LONGITUDE

Observation - There's no visible relationship between age and geographical coordinates (latitude and longitude). The points are dispersed uniformly.

Interpretation - Age is not influenced by or correlated with the location of the accident, which aligns with expectations since age should not directly impact where an accident occurs.

POSTED_SPEED_LIMIT vs. Other Features

POSTED_SPEED_LIMIT vs. NUM_UNITS

Observation - There's a slight pattern where accidents with higher speed limits involve slightly fewer vehicles, but the relationship is weak.

Interpretation - While there might be a mild trend that higher speed limits involve fewer vehicles, it is not a strong correlation. This could suggest that speed and the number of vehicles involved operate relatively independently.

POSTED_SPEED_LIMIT vs. LATITUDE/LONGITUDE

Observation - No clear patterns are visible, indicating that speed limits do not vary significantly by location within the region covered by the dataset.

Interpretation - This suggests that the speed limit is fairly consistent across different geographical areas in the dataset.

NUM_UNITS vs. Other Features

NUM_UNITS vs. LATITUDE/LONGITUDE

Observation - Similar to the other plots, there's no distinct relationship between the number of units involved in an accident and the geographical coordinates.

Interpretation - The number of units involved in accidents does not vary significantly by location, suggesting that accident severity or scale (in terms of units involved) is not location-dependent within the dataset.

LATITUDE vs. LONGITUDE Observation - The scatter plot between LATITUDE and LONGITUDE shows a strong linear relationship, which is expected since they represent geographic coordinates. The linearity reflects the physical layout of the area where the data was collected.

Interpretation - The strong correlation between latitude and longitude reinforces that these features are related to the same underlying factor (location). For modeling purposes, these may be combined or used in location-based analysis.

Conclusion Independence of Features - Most numerical features, such as age, speed limit, and number of units, show little to no correlation with each other. This suggests that these features can independently contribute to the predictive power of your model.

Geographical Coordinates - The strong correlation between latitude and longitude is expected, but they do not show any relationship with other numerical features like age or speed limit.

Weak Relationships - The scatter plots suggest that the numerical features may not be strongly interrelated, which is beneficial for avoiding multicollinearity in the model.

Exploring Relationships Between Categorical Variables and the Target Variable

```
In [60]: # Cross-tabulation example: VEHICLE_TYPE vs INJURY_CLASSIFICATION
cross_tab = pd.crosstab(merged_df_final['VEHICLE_TYPE'], merged_df_final['INJURY_CLASSIFICATION'])
print("Cross-Tabulation between Vehicle Type and Injury Classification:\n",
```

Cross-Tabulation between Vehicle Type and Injury Classification:

INJURY_CLASSIFICATION	FATAL	INCAPACITATING INJURY \
VEHICLE_TYPE		
3-WHEELED MOTORCYCLE (2 REAR WHEELS)	0	7
ALL-TERRAIN VEHICLE (ATV)	0	24
AUTOCYCLE	1	3
BUS OVER 15 PASS.	19	271
BUS UP TO 15 PASS.	1	24
FARM EQUIPMENT	0	1
MOPED OR MOTORIZED BICYCLE	4	65
MOTOR DRIVEN CYCLE	2	40
MOTORCYCLE (OVER 150CC)	60	543
OTHER	19	261
OTHER VEHICLE WITH TRAILER	1	18
PASSENGER	1215	21198
PICKUP	51	798
RECREATIONAL OFF-HIGHWAY VEHICLE (ROV)	0	4
SINGLE UNIT TRUCK WITH TRAILER	1	38
SNOWMOBILE	0	0
SPORT UTILITY VEHICLE (SUV)	157	3214
TRACTOR W/ SEMI-TRAILER	26	122
TRACTOR W/O SEMI-TRAILER	9	31
TRUCK - SINGLE UNIT	26	271
UNKNOWN/NA	86	1306
VAN/MINI-VAN	66	1249

INJURY_CLASSIFICATION	NO INDICATION OF INJURY \
VEHICLE_TYPE	
3-WHEELED MOTORCYCLE (2 REAR WHEELS)	68
ALL-TERRAIN VEHICLE (ATV)	225
AUTOCYCLE	1012
BUS OVER 15 PASS.	37132
BUS UP TO 15 PASS.	7988
FARM EQUIPMENT	119
MOPED OR MOTORIZED BICYCLE	494
MOTOR DRIVEN CYCLE	449
MOTORCYCLE (OVER 150CC)	4578
OTHER	30746
OTHER VEHICLE WITH TRAILER	3703
PASSENGER	1821315
PICKUP	84392
RECREATIONAL OFF-HIGHWAY VEHICLE (ROV)	27
SINGLE UNIT TRUCK WITH TRAILER	3581
SNOWMOBILE	9
SPORT UTILITY VEHICLE (SUV)	390581
TRACTOR W/ SEMI-TRAILER	22971
TRACTOR W/O SEMI-TRAILER	3334
TRUCK - SINGLE UNIT	48304
UNKNOWN/NA	202049
VAN/MINI-VAN	139828

INJURY_CLASSIFICATION	NONINCAPACITATING INJURY \
VEHICLE_TYPE	
3-WHEELED MOTORCYCLE (2 REAR WHEELS)	18
ALL-TERRAIN VEHICLE (ATV)	47
AUTOCYCLE	50
BUS OVER 15 PASS.	1795
BUS UP TO 15 PASS.	237
FARM EQUIPMENT	10
MOPED OR MOTORIZED BICYCLE	135
MOTOR DRIVEN CYCLE	105

MOTORCYCLE (OVER 150CC)	1178
OTHER	1249
OTHER VEHICLE WITH TRAILER	128
PASSENGER	111587
PICKUP	3991
RECREATIONAL OFF-HIGHWAY VEHICLE (ROV)	10
SINGLE UNIT TRUCK WITH TRAILER	107
SNOWMOBILE	0
SPORT UTILITY VEHICLE (SUV)	17803
TRACTOR W/ SEMI-TRAILER	728
TRACTOR W/O SEMI-TRAILER	138
TRUCK - SINGLE UNIT	1578
UNKNOWN/NA	6602
VAN/MINI-VAN	6820

INJURY_CLASSIFICATION	REPORTED, NOT EVIDENT
VEHICLE_TYPE	
3-WHEELED MOTORCYCLE (2 REAR WHEELS)	1
ALL-TERRAIN VEHICLE (ATV)	10
AUTOCYCLE	5
BUS OVER 15 PASS.	1250
BUS UP TO 15 PASS.	120
FARM EQUIPMENT	3
MOPED OR MOTORIZED BICYCLE	31
MOTOR DRIVEN CYCLE	22
MOTORCYCLE (OVER 150CC)	210
OTHER	677
OTHER VEHICLE WITH TRAILER	57
PASSENGER	61314
PICKUP	2435
RECREATIONAL OFF-HIGHWAY VEHICLE (ROV)	0
SINGLE UNIT TRUCK WITH TRAILER	58
SNOWMOBILE	0
SPORT UTILITY VEHICLE (SUV)	12243
TRACTOR W/ SEMI-TRAILER	487
TRACTOR W/O SEMI-TRAILER	89
TRUCK - SINGLE UNIT	1068
UNKNOWN/NA	3431
VAN/MINI-VAN	4412

Interpretation of the Cross-Tabulation between Vehicle Type and Injury Classification

Passenger Vehicles FATAL Injuries: Passenger vehicles are involved in a significantly higher number of fatal injuries (1,215 cases). **INCAPACITATING INJURY:** Passenger vehicles also lead in incapacitating injuries, with 21,198 cases. **NONINCAPACITATING INJURY:** Again, passenger vehicles are involved in the majority of non-incapacitating injuries (49,444 cases).

Interpretation: Passenger vehicles are the most common vehicle type involved in accidents, which explains their high numbers across all injury classifications. Their involvement in severe injuries (both fatal and incapacitating) highlights the importance of focusing on safety measures for passenger vehicles.

Sport Utility Vehicles (SUVs) FATAL Injuries: SUVs are involved in 157 fatal injuries.

INCAPACITATING INJURY: There are 3,214 cases of incapacitating injuries involving SUVs.

NONINCAPACITATING INJURY: SUVs are involved in 6,409 non-incapacitating injuries.

Interpretation: SUVs also show high numbers across all injury classifications, indicating that, similar to passenger vehicles, SUVs are commonly involved in accidents that result in injuries. This might be due to their popularity and prevalence on the roads.

Motorcycles (Over 150cc) FATAL Injuries: Motorcycles are associated with a significant number of fatal injuries (60 cases).

INCAPACITATING INJURY: There are 543 cases of incapacitating injuries involving motorcycles.

NONINCAPACITATING INJURY: Motorcycles are involved in 918 non-incapacitating injuries.

Interpretation: Motorcycles, although less common than passenger vehicles, have a notably high rate of severe injuries relative to their numbers, especially fatal and incapacitating injuries. This suggests that accidents involving motorcycles are more likely to result in serious injury, likely due to the lack of protection for riders.

Trucks and Commercial Vehicles FATAL Injuries: Trucks and single-unit trucks are involved in several fatal injuries (e.g., 26 for tractor w/ semi-trailer, 51 for pickup trucks).

INCAPACITATING INJURY: These vehicles also show considerable numbers in incapacitating injuries. **NONINCAPACITATING INJURY:** These vehicles are involved in a fair number of non-incapacitating injuries, though lower than passenger vehicles and SUVs.

Interpretation: Commercial vehicles like trucks and tractors, while not as frequently involved in accidents as passenger vehicles, still contribute to a significant number of severe injuries. This highlights the potential risks associated with larger vehicles.

Unknown/NA Vehicle Types FATAL Injuries: Unknown/NA vehicle types are involved in 86 fatal injuries.

INCAPACITATING INJURY: There are 1,306 cases of incapacitating injuries involving these vehicles.

NONINCAPACITATING INJURY: These vehicles are involved in 3,431 non-incapacitating injuries.

Interpretation - The "Unknown/NA" category represents vehicles that were either not properly identified or categorized. The relatively high numbers in this category suggest that data quality or vehicle identification might be an issue, and improving this could lead to better insights.

Conclusion Passenger Vehicles and SUVs: These are the most frequently involved in accidents, leading in all injury classifications. This is likely due to their prevalence on the road.

Motorcycles: Despite lower overall numbers, motorcycles are disproportionately involved in severe injuries, indicating a higher risk associated with motorcycle accidents.

Commercial Vehicles: Trucks and similar vehicles contribute significantly to severe injuries, underscoring the risks associated with larger vehicles.

Feature Exploration

Explore the Distribution of Numerical Features

In [61]:

```
# List of numerical features to explore
numerical_features = ['AGE', 'POSTED_SPEED_LIMIT', 'NUM_UNITS', 'LATITUDE',

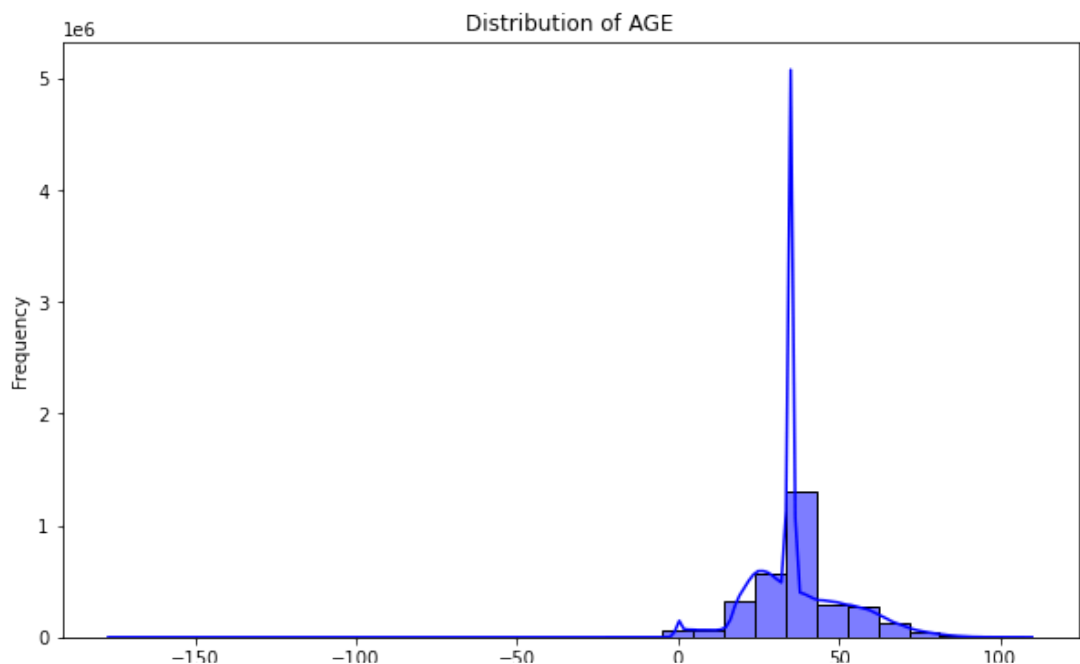
# Base path where images will be saved
save_base_path = r'C:\Users\MNJORGE16\Desktop\Moringa\phase_3\project__pha

# Plot histograms for numerical features and save them
for feature in numerical_features:
    plt.figure(figsize=(10, 6))
    sns.histplot(merged_df_final[feature], bins=30, kde=True, color='blue')
    plt.title(f'Distribution of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Frequency')

    # Define the save path for each plot
    save_path = f'{save_base_path}\\distribution_{feature}.png'

    # Save the figure as a PNG file
    plt.savefig(save_path)

    # Show the plot
    plt.show()
```



Interpretation of the distributions of the numerical features

Distribution of AGE

Observation - The AGE distribution is heavily skewed, with a sharp peak around the 30-40 age range. There are also extreme outliers with negative and very high ages, which are likely data entry errors.

Interpretation - The peak around 30-40 years suggests that this age group is the most frequently involved in traffic accidents. The negative and extremely high values indicate data issues that should be addressed (e.g., by removing or correcting these outliers). For modeling, age is expected to be a significant factor in predicting injury severity, but the data quality needs to be improved for accurate predictions.

Distribution of LATITUDE

Observation - The LATITUDE values are clustered within a narrow range, which aligns with the geographical area covered by the dataset. The distribution shows that most of the data points are concentrated in a specific latitude range.

Interpretation - This indicates that the accidents are occurring within a specific geographical area, likely corresponding to the city or region being studied. Since latitude alone might not directly impact injury classification, it could be combined with longitude or used in geospatial analyses.

Distribution of LONGITUDE

Observation - Similar to latitude, the LONGITUDE values are tightly clustered within a narrow range, with the vast majority of values within a specific interval.

Interpretation - The longitude distribution supports the finding that the dataset is geographically concentrated in a particular region. Like latitude, longitude may not directly influence injury severity but could be useful in combination with other features or for location-based analysis.

Distribution of NUM_UNITS

Observation - The NUM_UNITS (number of units involved in the crash) distribution is heavily skewed to the left, with the majority of accidents involving 2-3 units. There are a few outliers with a higher number of units.

Interpretation - Most accidents involve a small number of units, typically 2-3 vehicles. The presence of outliers suggests that some accidents involve significantly more vehicles, which could be associated with more complex scenarios or higher severity, but these cases are rare.

Distribution of POSTED_SPEED_LIMIT

Observation - The POSTED_SPEED_LIMIT distribution shows a sharp peak around 20-40 mph, which is typical for urban areas. There are outliers at both the low and high ends of the speed limit range.

Interpretation - The concentration around 20-40 mph suggests that most accidents occur in urban settings where these speed limits are common. The outliers at lower and higher speed limits might correspond to rural or highway areas. The posted speed limit is likely an important factor in determining injury severity, especially when combined with other features like vehicle type or maneuver.

Conclusion

Age - This is a key feature with a significant peak in the 30-40 age range, but data cleaning is needed due to outliers.

Latitude and Longitude - Both are tightly clustered, indicating a specific geographical focus. They might not be directly predictive of injury classification but could be used in combination for location-based analysis.

Number of Units - Most accidents involve a small number of units, which could influence injury severity predictions.

Examine Relationships Between Numerical Features and the Target Variable

```
In [62]: # Plot box plots for numerical features against the target variable

# List of numerical features to explore
numerical_features = ['AGE', 'POSTED_SPEED_LIMIT', 'NUM_UNITS', 'LATITUDE',

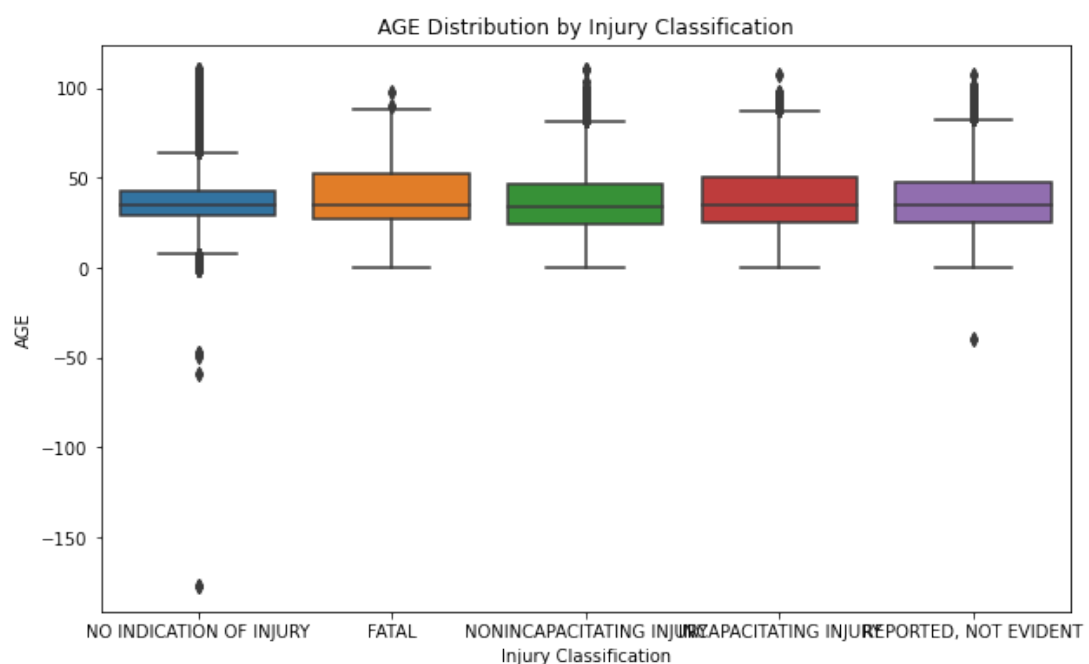
# Base path where images will be saved
save_base_path = r'C:\Users\MNJORGE16\Desktop\Moringa\phase_3\project__pha

# Plot box plots for numerical features against the target variable and save
for feature in numerical_features:
    plt.figure(figsize=(10, 6))
    sns.boxplot(x='INJURY_CLASSIFICATION', y=feature, data=merged_df_final)
    plt.title(f'{feature} Distribution by Injury Classification')
    plt.xlabel('Injury Classification')
    plt.ylabel(feature)

# Define the save path for each plot
save_path = f'{save_base_path}\\boxplot_{feature}_vs_injury_classification

# Save the figure as a PNG file
plt.savefig(save_path)

# Show the plot
plt.show()
```



Boxplot Output Interpretation

AGE Distribution by Injury Classification Observation - The AGE distribution across different injury classifications is relatively similar, with the median age being around 30-40 years old for all injury types. There are outliers, with some negative and extremely high values, which might indicate data entry errors.

Interpretation - The similarity in age distribution across injury classifications suggests that age alone might not be a strong predictor of injury severity. However, the presence of outliers, especially negative values, indicates a need for further data cleaning or handling of these erroneous entries.

LATITUDE Distribution by Injury Classification Observation - The LATITUDE values are clustered tightly around a narrow range, likely corresponding to the geographical region covered by the dataset. There are outliers with very low values, which might be erroneous or indicate locations outside the expected range.

Interpretation - The tight clustering of latitude values suggests that the accidents occur within a specific geographical area. The outliers could represent data errors or unusual cases that might need to be handled separately.

LONGITUDE Distribution by Injury Classification

Observation - Similar to latitude, the LONGITUDE values are clustered within a specific range, with outliers that have very low or negative values. Interpretation: The longitude data shows a similar pattern to latitude, with most data points concentrated in a specific geographical region. The outliers here also suggest potential data entry errors or unusual cases.

NUM_UNITS Distribution by Injury Classification

Observation - The NUM_UNITS (number of units involved in the crash) generally ranges between 2 and 3 for most injury classifications, with higher numbers being less common. Outliers with higher numbers of units involved are present, particularly in the "NO INDICATION OF INJURY" category.

Interpretation - The distribution suggests that most accidents involve a small number of units (likely vehicles). Higher numbers of units involved don't necessarily correlate with more severe injuries, as they also appear in non-injury cases. This might indicate that the number of vehicles involved isn't a straightforward predictor of injury severity but could be a contributing factor when combined with other variables.

POSTED_SPEED_LIMIT Distribution by Injury Classification Observation - The POSTED_SPEED_LIMIT values show a wide range, with most data points clustering around typical urban speed limits (20-40 mph). Outliers exist at both ends of the spectrum, particularly in the "NO INDICATION OF INJURY" category.

Interpretation - The speed limit at the site of the crash varies widely but is most often in the 20-40 mph range, which is common in urban settings. The presence of outliers suggests that very high or very low-speed limits are less common but do exist. This distribution might indicate that the posted speed limit alone is not a strong predictor of injury severity, but it could be an important factor when considered alongside other variables like road conditions or vehicle type.

Explore Categorical Features and Their Relationship with the Target Variable

```
In [63]: # List of categorical features to explore
categorical_features = ['VEHICLE_TYPE', 'TRAVEL_DIRECTION', 'MANEUVER', 'SE

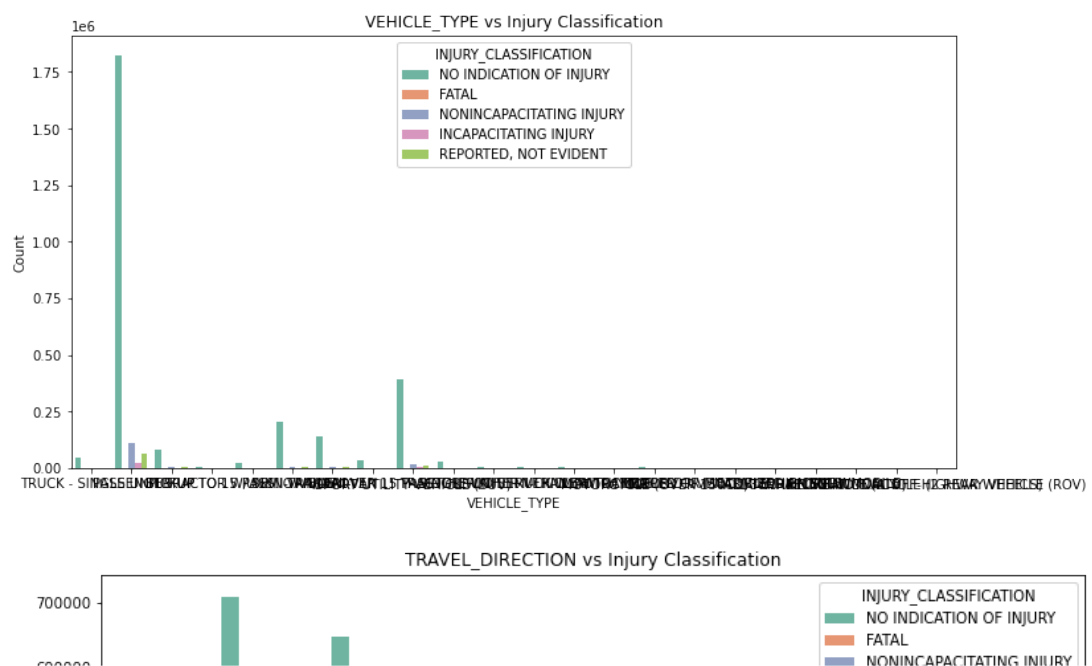
# Base path where images will be saved
save_base_path = r'C:\Users\MNJORGE16\Desktop\Moringa\phase_3\project__pha

# Plot bar plots for categorical features and save them
for feature in categorical_features:
    plt.figure(figsize=(12, 6))
    sns.countplot(x=feature, hue='INJURY_CLASSIFICATION', data=merged_df_fi
    plt.title(f'{feature} vs Injury Classification')
    plt.xlabel(feature)
    plt.ylabel('Count')

# Define the save path for each plot
save_path = f'{save_base_path}\\barplot_{feature}_vs_injury_classificat

# Save the figure as a PNG file
plt.savefig(save_path)

# Show the plot
plt.show()
```



Barplots Interpretation

MANEUVER vs Injury Classification

Observation - The majority of maneuvers are concentrated around the "STRAIGHT AHEAD" maneuver, with a very high count in the "NO INDICATION OF INJURY" category. Other maneuvers have significantly lower counts. Interpretation - This suggests that most accidents occur while vehicles are moving straight ahead, and these incidents are often

non-injurious. This could indicate that straight driving is common, but when incidents occur, they tend to be less severe. However, certain maneuvers with lower frequencies might be associated with more severe injuries.

SAFETY_EQUIPMENT vs Injury Classification

Observation - The vast majority of individuals were using "SAFETY BELT" equipment, predominantly resulting in "NO INDICATION OF INJURY". Other categories such as "NONE" have much lower frequencies but show higher incidences of injuries.

Interpretation - The data indicates that the use of safety equipment, especially safety belts, is strongly associated with a lower risk of injury. This supports the effectiveness of safety belts in reducing injury severity in traffic accidents.

SEX vs Injury Classification

Observation - There is a higher number of male ("M") participants in the data, with the majority having "NO INDICATION OF INJURY". Females ("F") also show a similar pattern but with fewer occurrences.

Interpretation - This suggests that more males are involved in accidents than females, but the injury distribution between sexes seems relatively similar, with the majority of both males and females not sustaining injuries in these incidents.

TRAVEL_DIRECTION vs Injury Classification

Observation - The most common travel directions are "N", "S", "E", and "W", with "NO INDICATION OF INJURY" being the most frequent classification for all directions. The "UNKNOWN" direction also appears but has fewer entries and a notable number of injuries.

Interpretation - The direction of travel appears to have little impact on injury severity, as most incidents across all directions result in no injuries. However, the "UNKNOWN" direction might be associated with less typical or more severe circumstances leading to injuries.

VEHICLE_TYPE vs Injury Classification Observation - The "PASSENGER" vehicle type dominates the data, with a very high count under "NO INDICATION OF INJURY". Other vehicle types, like "TRUCK" or "SUV", have significantly lower frequencies.

Interpretation - Most accidents involve passenger vehicles, which are more common on the roads, and the majority of these incidents do not result in injury. However, the data may

Data Preparation

Feature Selection

Review Data Understanding Insights

Objective - Reassess the key insights from the data understanding phase

```
In [64]: # List of initial features after data understanding
initial_features = ['AGE', 'POSTED_SPEED_LIMIT', 'NUM_UNITS', 'LATITUDE', '
                'TRAVEL_DIRECTION', 'MANEUVER', 'SEX', 'SAFETY_EQUIPMEN

# Target variable
target = 'INJURY_CLASSIFICATION'

# Review correlation matrix and scatter plots (already done in data underst
# No additional code needed here; just use insights from the previous steps
```

Remove Irrelevant Features

Objective - Drop features that are irrelevant or have little predictive power.

```
In [65]: # Dropping irrelevant or redundant features identified earlier
features_to_drop = ['LATITUDE', 'LONGITUDE']
df_reduced = merged_df_final.drop(columns=features_to_drop)
print("Remaining features after dropping irrelevant ones:", df_reduced.colu
```

```
Remaining features after dropping irrelevant ones: Index(['CRASH_RECORD_I
D', 'CRASH_DATE_x', 'POSTED_SPEED_LIMIT',
                'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION', 'WEATHER_CONDITION',
                'LIGHTING_CONDITION', 'FIRST_CRASH_TYPE', 'TRAFFICWAY_TYPE',
                'ALIGNMENT', 'ROADWAY_SURFACE_COND', 'ROAD_DEFECT', 'REPORT_TYPE',
                'CRASH_TYPE', 'DAMAGE', 'DATE_POLICE_NOTIFIED',
                'PRIM_CONTRIBUTORY_CAUSE', 'SEC_CONTRIBUTORY_CAUSE', 'STREET_NO',
                'STREET_DIRECTION', 'STREET_NAME', 'BEAT_OF_OCCURRENCE', 'NUM_UNIT
S',
                'MOST_SEVERE_INJURY', 'INJURIES_TOTAL', 'INJURIES_FATAL',
                'INJURIES_INCAPACITATING', 'INJURIES_NON_INCAPACITATING',
                'INJURIES_REPORTED_NOT_EVIDENT', 'INJURIES_NO_INDICATION',
                'INJURIES_UNKNOWN', 'CRASH_HOUR', 'CRASH_DAY_OF_WEEK', 'CRASH_MONT
H',
                'LOCATION', 'PERSON_ID', 'PERSON_TYPE', 'VEHICLE_ID_x', 'CRASH_DAT
E_y',
                'SEX', 'AGE', 'DRIVERS_LICENSE_STATE', 'DRIVERS_LICENSE_CLASS',
                'SAFETY_EQUIPMENT', 'AIRBAG_DEPLOYED', 'EJECTION',
                'INJURY_CLASSIFICATION', 'DRIVER_ACTION', 'DRIVER_VISION',
                'PHYSICAL_CONDITION', 'BAC_RESULT', 'CRASH_UNIT_ID', 'CRASH_DATE',
                'UNIT_NO', 'UNIT_TYPE', 'VEHICLE_ID_y', 'MAKE', 'MODEL',
                'LIC_PLATE_STATE', 'VEHICLE_YEAR', 'VEHICLE_DEFECT', 'VEHICLE_TYP
E',
                'VEHICLE_USE', 'TRAVEL_DIRECTION', 'MANEUVER', 'OCCUPANT_CNT',
                'FIRST_CONTACT_POINT'],
                dtype='object')
```

Address Multicollinearity


Objective - Handling multicollinearity, important for logistic regression.

```
In [66]: # Calculate the correlation matrix
corr_matrix = df_reduced.corr()

# Plot the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix After Dropping Features')

# Save the heatmap to the specified folder
output_path = r'C:\Users\MNJ0R0GE16\Desktop\Moringa\phase_3\project__phase3
plt.savefig(output_path)

# Display the plot
plt.show()
```



Interpretation of the Correlation Matrix Heatmap

Low Correlation Between Most Features

The majority of the features have low correlation values (close to 0), indicating that they do not have strong linear relationships with each other. This is beneficial for both logistic regression and decision trees because.

Logistic Regression Low correlation reduces the risk of multicollinearity, which can lead to unreliable estimates of coefficients.

Decision Trees The algorithm is robust to multicollinearity, but having uncorrelated features ensures that each feature contributes unique information to the model.

Highly Correlated Features

Injuries Features The features related to different types of injuries (e.g., INJURIES_TOTAL, INJURIES_FATAL, INCAPACITATING, etc.) show strong correlations with each other.

Action - For logistic regression, you might consider dropping or combining these highly correlated features to avoid multicollinearity. For decision trees, you might retain them, as the model can handle correlated features well.

Vehicle IDs Features like VEHICLE_ID_x and VEHICLE_ID_y are also highly correlated.

Action - These are likely identifiers or categorical variables that may not be necessary for modeling. Consider dropping them unless they provide meaningful insights.

Conclusion

Feature Selection The heatmap suggests focusing on features that are less correlated with each other, which can help in building more stable models.

Injury Features Given that many injury-related features are highly correlated, you might choose the most representative one(s) or create composite scores (e.g., summing or averaging certain features) for logistic regression.

Geographical Features LATITUDE and LONGITUDE were dropped due to low relevance or high correlation with each other. This decision appears justified given the project

Type *Markdown* and LaTeX: α^2

In [67]:

```
#Review the correlation matrix
corr_matrix = df_reduced.corr().abs() # Use absolute values to consider po

#Identify features with high correlation
# Setting a threshold of 0.8 as an example for high correlation
high_corr_pairs = np.where(corr_matrix > 0.8)
high_corr_pairs = [(corr_matrix.index[x], corr_matrix.columns[y])
                    for x, y in zip(*high_corr_pairs)
                    if x != y and x < y]

print("Highly correlated feature pairs:")
for pair in high_corr_pairs:
    print(pair)

# Drop redundant features based on correlation and domain knowledge
# Dropping VEHICLE_ID_x and VEHICLE_ID_y as they are identifiers and not us
# Drop additional features identified from correlation matrix review
features_to_drop = [
    'NUM_UNITS', 'INJURIES_TOTAL', 'INJURIES_FATAL', 'INJURIES_INCAPACITATI
    'INJURIES_NON_INCAPACITATING', 'INJURIES_REPORTED_NOT_EVIDENT', 'INJURI
    'INJURIES_UNKNOWN', 'OCCUPANT_CNT']

df_final = df_reduced.drop(columns=features_to_drop)
print("Final feature set after removing redundant features:")
print(df_final.columns)

# Save the final refined dataset
df_final.to_csv(r'C:\Users\MNJOROG16\Desktop\Moringa\phase_3\project__phas
```

```
Highly correlated feature pairs:
('VEHICLE_ID_x', 'CRASH_UNIT_ID')
('VEHICLE_ID_x', 'VEHICLE_ID_y')
('CRASH_UNIT_ID', 'VEHICLE_ID_y')
Final feature set after removing redundant features:
Index(['CRASH_RECORD_ID', 'CRASH_DATE_x', 'POSTED_SPEED_LIMIT',
      'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION', 'WEATHER_CONDITION',
      'LIGHTING_CONDITION', 'FIRST_CRASH_TYPE', 'TRAFFICWAY_TYPE',
      'ALIGNMENT', 'ROADWAY_SURFACE_COND', 'ROAD_DEFECT', 'REPORT_TYPE',
      'CRASH_TYPE', 'DAMAGE', 'DATE_POLICE_NOTIFIED',
      'PRIM_CONTRIBUTORY_CAUSE', 'SEC_CONTRIBUTORY_CAUSE', 'STREET_NO',
      'STREET_DIRECTION', 'STREET_NAME', 'BEAT_OF_OCCURRENCE',
      'MOST_SEVERE_INJURY', 'CRASH_HOUR', 'CRASH_DAY_OF_WEEK', 'CRASH_MO
NTH',
      'LOCATION', 'PERSON_ID', 'PERSON_TYPE', 'VEHICLE_ID_x', 'CRASH_DAT
E_y',
      'SEX', 'AGE', 'DRIVERS_LICENSE_STATE', 'DRIVERS_LICENSE_CLASS',
      'SAFETY_EQUIPMENT', 'AIRBAG_DEPLOYED', 'EJECTION',
      'INJURY_CLASSIFICATION', 'DRIVER_ACTION', 'DRIVER_VISION',
      'PHYSICAL_CONDITION', 'BAC_RESULT', 'CRASH_UNIT_ID', 'CRASH_DATE',
      'UNIT_NO', 'UNIT_TYPE', 'VEHICLE_ID_y', 'MAKE', 'MODEL',
      'LIC_PLATE_STATE', 'VEHICLE_YEAR', 'VEHICLE_DEFECT', 'VEHICLE_TYP
E',
      'VEHICLE_USE', 'TRAVEL_DIRECTION', 'MANEUVER', 'FIRST_CONTACT_POIN
T'],
      dtype='object')
```

Handle Remaining Missing Values

Missing values were handled during data understanding

Feature Engineering

Create New Features

New features These features help the model understand temporal patterns, such as whether certain times of the day or week are associated with higher injury severity.

Interaction features To create meaningful interactions that can enhance the predictive power of the models

Binning Features Binning helps to simplify the relationship between age and injury severity by grouping ages into broader categories.

```
In [68]: # Create time-based features
df_final['CRASH_HOUR'] = pd.to_datetime(df_final['CRASH_DATE_x']).dt.hour
df_final['CRASH_DAY_OF_WEEK'] = pd.to_datetime(df_final['CRASH_DATE_x']).dt.dayofweek
df_final['CRASH_MONTH'] = pd.to_datetime(df_final['CRASH_DATE_x']).dt.month

print("Created time-based features: CRASH_HOUR, CRASH_DAY_OF_WEEK, CRASH_MONTH")
```

Created time-based features: CRASH_HOUR, CRASH_DAY_OF_WEEK, CRASH_MONTH

```
In [69]: # Create an interaction feature between POSTED_SPEED_LIMIT and VEHICLE_TYPE

#The type of vehicle involved in a crash combined with the speed limit can
# For example, crashes involving motorcycles at high speeds might result in
df_final['SPEED_VEHICLE_TYPE'] = df_final['POSTED_SPEED_LIMIT'] * df_final['VEHICLE_TYPE']

print("Created interaction feature: SPEED_VEHICLE_TYPE")
```

Created interaction feature: SPEED_VEHICLE_TYPE

```
In [70]: # Create an interaction feature between POSTED_SPEED_LIMIT and VEHICLE_TYPE
#The combination of weather and lighting conditions can significantly impact
df_final['SPEED_VEHICLE_TYPE'] = df_final['POSTED_SPEED_LIMIT'] * df_final[
print("Created interaction feature: SPEED_VEHICLE_TYPE")
```

Created interaction feature: SPEED_VEHICLE_TYPE

```
In [71]: # Binning AGE into categories
bins = [0, 18, 30, 50, 70, 100]
labels = ['Youth', 'Young Adult', 'Adult', 'Senior', 'Elder']
df_final['AGE_BINNED'] = pd.cut(df_final['AGE'], bins=bins, labels=labels,
print("Binned AGE into categories: Youth, Young Adult, Adult, Senior, Elder")
```

Binned AGE into categories: Youth, Young Adult, Adult, Senior, Elder

####

Feature Encoding

Identify Categorical Features

TRAFFIC_CONTROL_DEVICE

DEVICE_CONDITION

WEATHER_CONDITION

LIGHTING_CONDITION

FIRST_CRASH_TYPE

TRAFFICWAY_TYPE

ROADWAY_SURFACE_COND

ROAD_DEFECT

REPORT_TYPE

CRASH_TYPE

MOST_SEVERE_INJURY

PERSON_TYPE

SEX

DRIVERS_LICENSE_STATE

DRIVERS_LICENSE_CLASS

SAFETY_EQUIPMENT

AIRBAG_DEPLOYED

EJECTION

INJURY_CLASSIFICATION

DRIVER_ACTION

DRIVER_VISION

PHYSICAL_CONDITION

VEHICLE_TYPE

VEHICLE_USE

TRAVEL_DIRECTION

MANEUVER

One-Hot Encoding

One-Hot Encoding step is crucial because it ensures that all categorical data is in a format that can be utilized by logistic regression and decision tree models, ultimately contributing to more accurate predictions of injury severity.

In [72]:

```
# List of categorical features to encode
categorical_columns = [
    'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION', 'WEATHER_CONDITION',
    'LIGHTING_CONDITION', 'FIRST_CRASH_TYPE', 'TRAFFICWAY_TYPE',
    'ROADWAY_SURFACE_COND', 'ROAD_DEFECT', 'REPORT_TYPE',
    'CRASH_TYPE', 'MOST_SEVERE_INJURY', 'PERSON_TYPE',
    'SEX', 'DRIVERS_LICENSE_STATE', 'DRIVERS_LICENSE_CLASS',
    'SAFETY_EQUIPMENT', 'AIRBAG_DEPLOYED', 'EJECTION',
    'INJURY_CLASSIFICATION', 'DRIVER_ACTION', 'DRIVER_VISION',
    'PHYSICAL_CONDITION', 'VEHICLE_TYPE', 'VEHICLE_USE',
    'TRAVEL_DIRECTION', 'MANEUVER'
]

# Apply one-hot encoding to categorical columns
df_encoded = pd.get_dummies(df_final, columns=categorical_columns, drop_fir
```

In [73]:

```
# Display the shape and columns of the resulting DataFrame
print("Data shape after one-hot encoding:", df_encoded.shape)
```

Data shape after one-hot encoding: (3076376, 749)

```
In [76]: print("Encoded columns:")
print(df_encoded.columns)
```

```
Encoded columns:
Index(['CRASH_RECORD_ID', 'CRASH_DATE_x', 'POSTED_SPEED_LIMIT', 'ALIGNMEN
T',
      'DAMAGE', 'DATE_POLICE_NOTIFIED', 'PRIM_CONTRIBUTORY_CAUSE',
      'SEC_CONTRIBUTORY_CAUSE', 'STREET_NO', 'STREET_DIRECTION',
      ...,
      'MANEUVER_SLOW/STOP - LOAD/UNLOAD', 'MANEUVER_SLOW/STOP - RIGHT TU
RN',
      'MANEUVER_SLOW/STOP IN TRAFFIC', 'MANEUVER_STARTING IN TRAFFIC',
      'MANEUVER_STRAIGHT AHEAD', 'MANEUVER_TURNING LEFT',
      'MANEUVER_TURNING ON RED', 'MANEUVER_TURNING RIGHT', 'MANEUVER_U-T
URN',
      'MANEUVER_UNKNOWN/NA'],
      dtype='object', length=749)
```

Feature Scaling

```
In [78]: # Check the columns available in df_encoded
print("Available columns in df_encoded:")
print(df_encoded.columns)
```

```
Available columns in df_encoded:
Index(['CRASH_RECORD_ID', 'CRASH_DATE_x', 'POSTED_SPEED_LIMIT', 'ALIGNMEN
T',
      'DAMAGE', 'DATE_POLICE_NOTIFIED', 'PRIM_CONTRIBUTORY_CAUSE',
      'SEC_CONTRIBUTORY_CAUSE', 'STREET_NO', 'STREET_DIRECTION',
      ...,
      'MANEUVER_SLOW/STOP - LOAD/UNLOAD', 'MANEUVER_SLOW/STOP - RIGHT TU
RN',
      'MANEUVER_SLOW/STOP IN TRAFFIC', 'MANEUVER_STARTING IN TRAFFIC',
      'MANEUVER_STRAIGHT AHEAD', 'MANEUVER_TURNING LEFT',
      'MANEUVER_TURNING ON RED', 'MANEUVER_TURNING RIGHT', 'MANEUVER_U-T
URN',
      'MANEUVER_UNKNOWN/NA'],
      dtype='object', length=749)
```

```
In [79]: # List of numerical features to scale (adjusted to match the available colu
numerical_columns = ['POSTED_SPEED_LIMIT', 'AGE']
```

```
In [81]: # Verify the availability of numerical columns before scaling
available_numerical_columns = [col for col in numerical_columns if col in c

if available_numerical_columns:
    # Apply scaling to available numerical features
    df_encoded[available_numerical_columns] = scaler.fit_transform(df_encoded[available_numerical_columns])

    print("Numerical features after scaling:")
    print(df_encoded[available_numerical_columns].head())
else:
    print("No numerical columns available for scaling.")
```

Numerical features after scaling:

	POSTED_SPEED_LIMIT	AGE
0	-0.678505	-0.754361
1	-0.678505	-0.754361
2	-0.678505	2.283221
3	-0.678505	2.283221
4	-0.678505	-0.133038

Split Data into Training and Testing Sets

Define the Feature Matrix (X) and Target Variable (y)

Objective - Separate the dataset into the input features (X) and the target variable (y).

Drop the target variable from the dataset to create X and set y to the target column (INJURY_CLASSIFICATION).

Split the Data

Objective - Divide the dataset into training and testing sets to evaluate the performance of the models

Use `train_test_split` to create the training and testing datasets.

```
In [85]: # Identify the one-hot encoded columns related to INJURY_CLASSIFICATION
target_columns = [
    'INJURY_CLASSIFICATION_INCAPACITATING INJURY',
    'INJURY_CLASSIFICATION_NO INDICATION OF INJURY',
    'INJURY_CLASSIFICATION_NONINCAPACITATING INJURY',
    'INJURY_CLASSIFICATION_REPORTED, NOT EVIDENT'
]

# Combine the one-hot encoded columns back into a single categorical column
df_encoded['INJURY_CLASSIFICATION'] = df_encoded[target_columns].idxmax(axis=1)

# Drop the one-hot encoded columns since they are merged back
df_encoded = df_encoded.drop(columns=target_columns)

# Define the feature matrix (X) and target variable (y)
y = df_encoded['INJURY_CLASSIFICATION']
X = df_encoded.drop(columns=['INJURY_CLASSIFICATION']) # Dropping the target
```

```
In [86]: from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

print("Training data shape (X_train):", X_train.shape)
print("Training labels shape (y_train):", y_train.shape)
print("Testing data shape (X_test):", X_test.shape)
print("Testing labels shape (y_test):", y_test.shape)
```

```
Training data shape (X_train): (2153463, 745)
Training labels shape (y_train): (2153463,)
Testing data shape (X_test): (922913, 745)
Testing labels shape (y_test): (922913,)
```


Modelling

Decision Tree

```
In [11]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder

# Identify high-cardinality categorical columns
categorical_columns = X.select_dtypes(include=['object']).columns
high_cardinality_columns = [col for col in categorical_columns if X[col].nunique() > 10]

# Apply Label encoding to high-cardinality columns
label_encoders = {}
for col in high_cardinality_columns:
    le = LabelEncoder()
    X[col] = le.fit_transform(X[col].astype(str))
    label_encoders[col] = le

# Apply one-hot encoding to the remaining categorical columns with low cardinality
low_cardinality_columns = [col for col in categorical_columns if col not in high_cardinality_columns]
X_encoded = pd.get_dummies(X, columns=low_cardinality_columns)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)

# Train a baseline Decision Tree model
dt_baseline = DecisionTreeClassifier(random_state=42)
dt_baseline.fit(X_train, y_train)

# Predict on the test set
y_pred = dt_baseline.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print("Baseline Decision Tree Model Accuracy:", accuracy)
print("\nClassification Report:\n", report)
```

Baseline Decision Tree Model Accuracy: 0.9685235769785451

Classification Report:

	precision	recall	f1-score	support
FATAL	0.88	0.90	0.89	518
INCAPACITATING INJURY	0.83	0.83	0.83	8924
NO INDICATION OF INJURY	0.98	0.98	0.98	840676
NONINCAPACITATING INJURY	0.83	0.82	0.82	46574
REPORTED, NOT EVIDENT	0.78	0.77	0.77	26221
accuracy			0.97	922913
macro avg	0.86	0.86	0.86	922913
weighted avg	0.97	0.97	0.97	922913

```

In [17]: #Evaluate Baseline Model
import os
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import matplotlib.pyplot as plt
import seaborn as sns

# Define the path where visualizations will be saved
visuals_path = r'C:\Users\MNJORGE16\Desktop\Moringa\phase_3\project_phase3\visuals'

# Ensure the directory exists
os.makedirs(visuals_path, exist_ok=True)

# Evaluate the baseline model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

print("Baseline Decision Tree Model Accuracy:", accuracy)
print("\nClassification Report:\n", report)

# Plotting Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.savefig(os.path.join(visuals_path, "baseline_confusion_matrix.png"))
plt.show()

# ROC and AUC
y_pred_proba = dt_baseline.predict_proba(X_test)
fpr, tpr, _ = roc_curve(y_test, y_pred_proba[:, 1], pos_label=dt_baseline.classes_[1])
roc_auc = auc(fpr, tpr)


plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color="blue", label=f"ROC curve (area = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], color="gray", linestyle="--")
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend(loc="lower right")
plt.savefig(os.path.join(visuals_path, "baseline_roc_curve.png"))
plt.show()


```

Baseline Decision Tree Model Accuracy: 0.9685235769785451

Classification Report:

	precision	recall	f1-score	support
FATAL	0.88	0.90	0.89	518
INCAPACITATING INJURY	0.83	0.83	0.83	8924
NO INDICATION OF INJURY	0.98	0.98	0.98	840676
NONINCAPACITATING INJURY	0.83	0.82	0.82	46574
REPORTED, NOT EVIDENT	0.78	0.77	0.77	26221
accuracy			0.97	922913
macro avg	0.86	0.86	0.86	922913
weighted avg	0.97	0.97	0.97	922913





```
In [23]: from sklearn.model_selection import train_test_split

# Reduce the dataset size using stratified sampling
X_small, _, y_small, _ = train_test_split(X, y, test_size=0.7, stratify=y,

# Print the size of the reduced dataset
print("Reduced dataset size:", X_small.shape, y_small.shape)
```

Reduced dataset size: (922912, 57) (922912,)